CompSci 401: Cloud Computing

Managing Distributed Systems

Prof. Ítalo Cunha

昆山杜克大学
DUKE KUNSHAN UNIVERSITY

Distributed driver's license management system

- One node per state
- Each node keeps information for drivers in that state
- Systems interoperate when cross-state information is needed

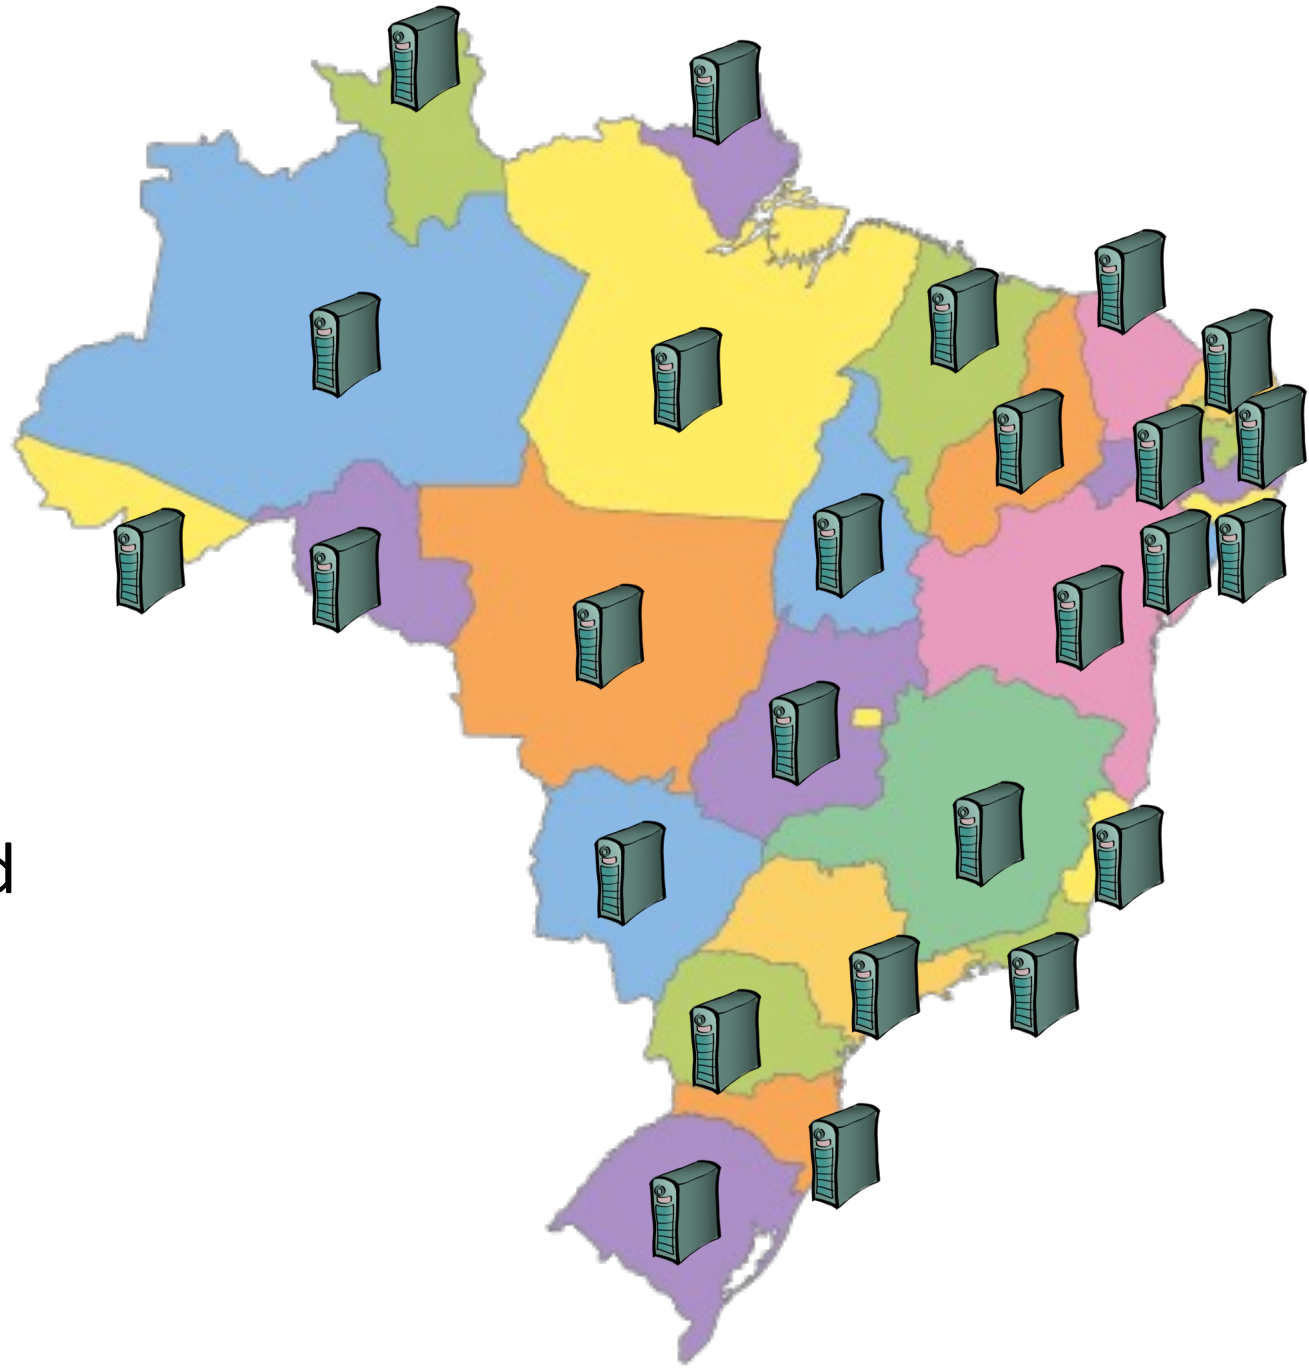

Distributed driver's license management system

- One node per state
- Each node keeps information for drivers in that state
- Systems interoperate when cross-state information is needed

- Efficient design if most accesses are local
  - For example, lower latency

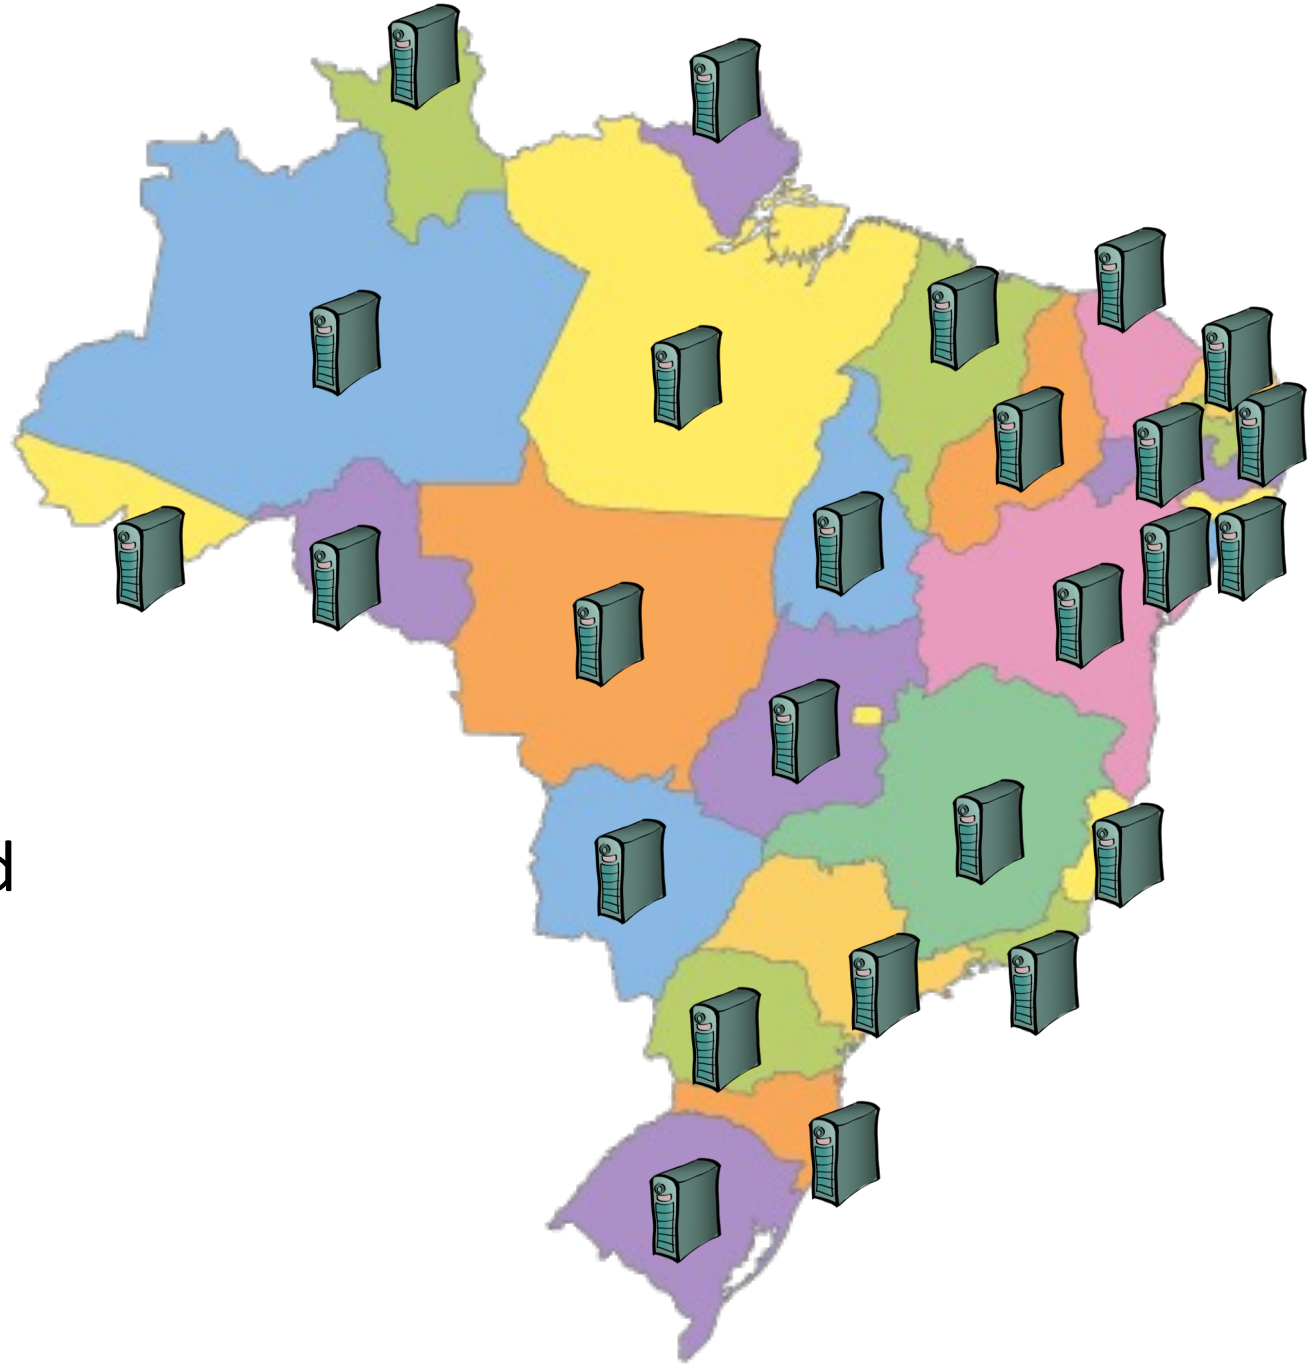

How to keep the system running?

- Local IT staff can troubleshoot
  - Reboot servers, restart nodes
  - Repair network links

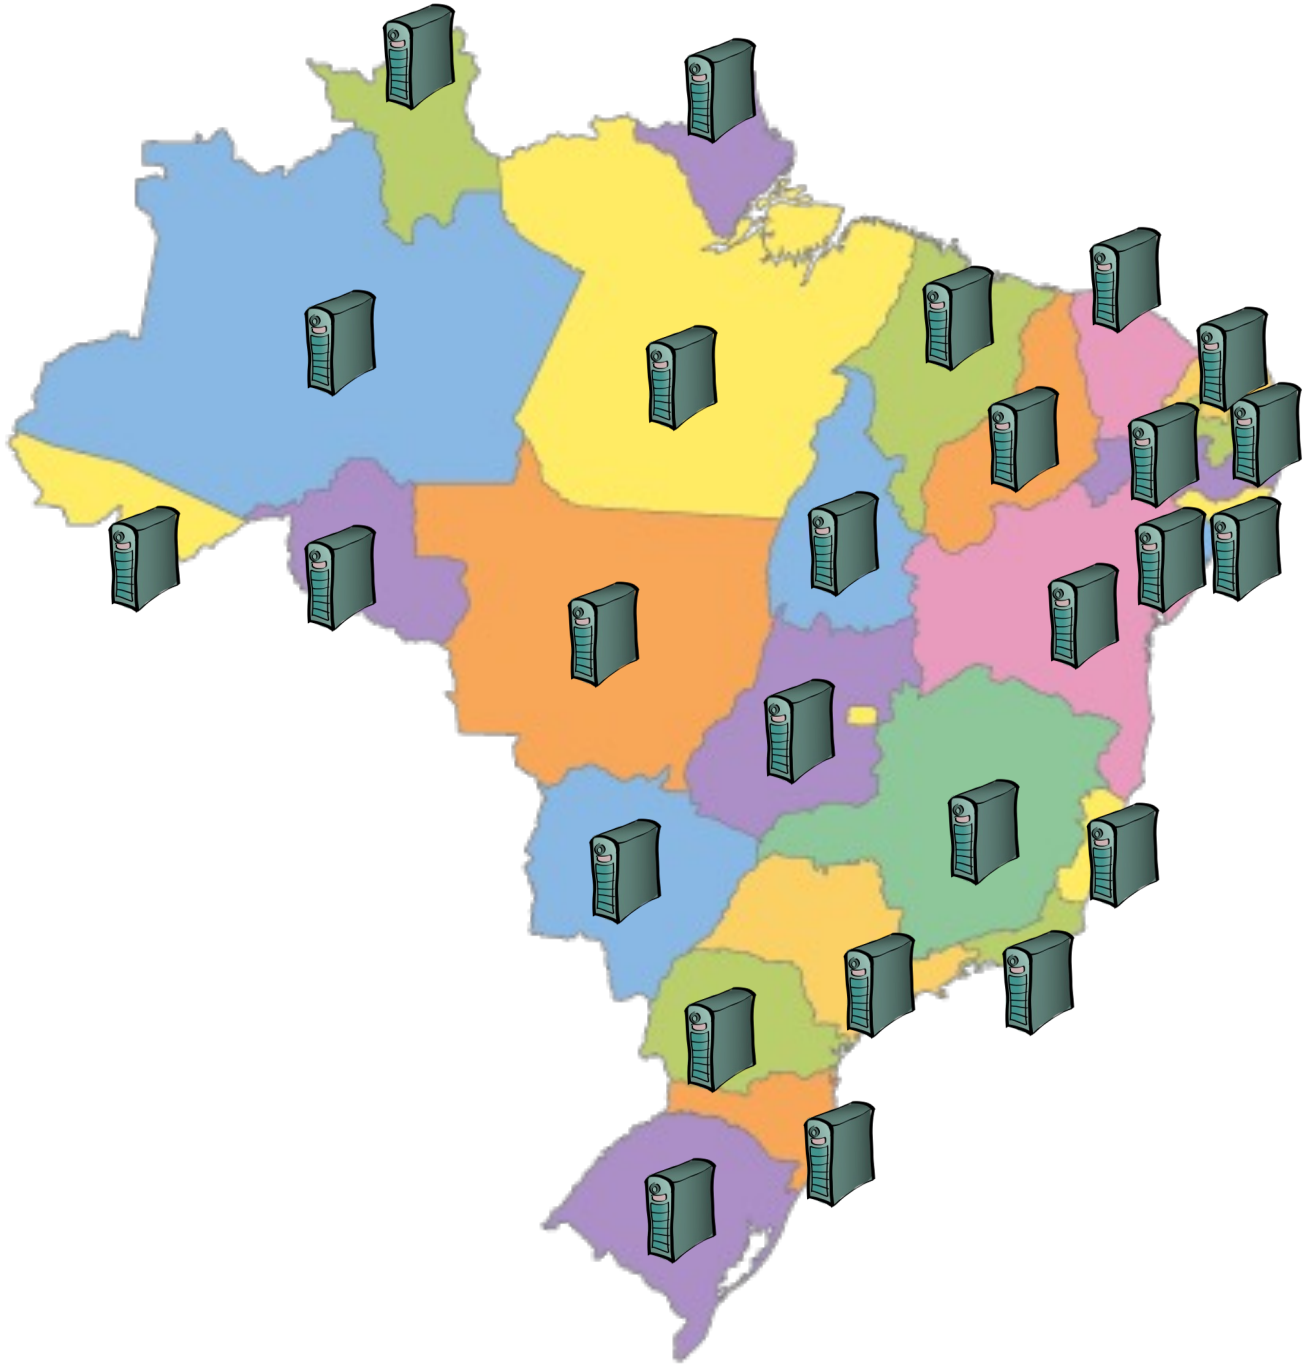

Monitoring a distributed system

- Periodically probe nodes
  - What frequency?
  - Liveliness vs overhead
- How to probe a node
  - What to check?
  - Thoroughness vs complexity

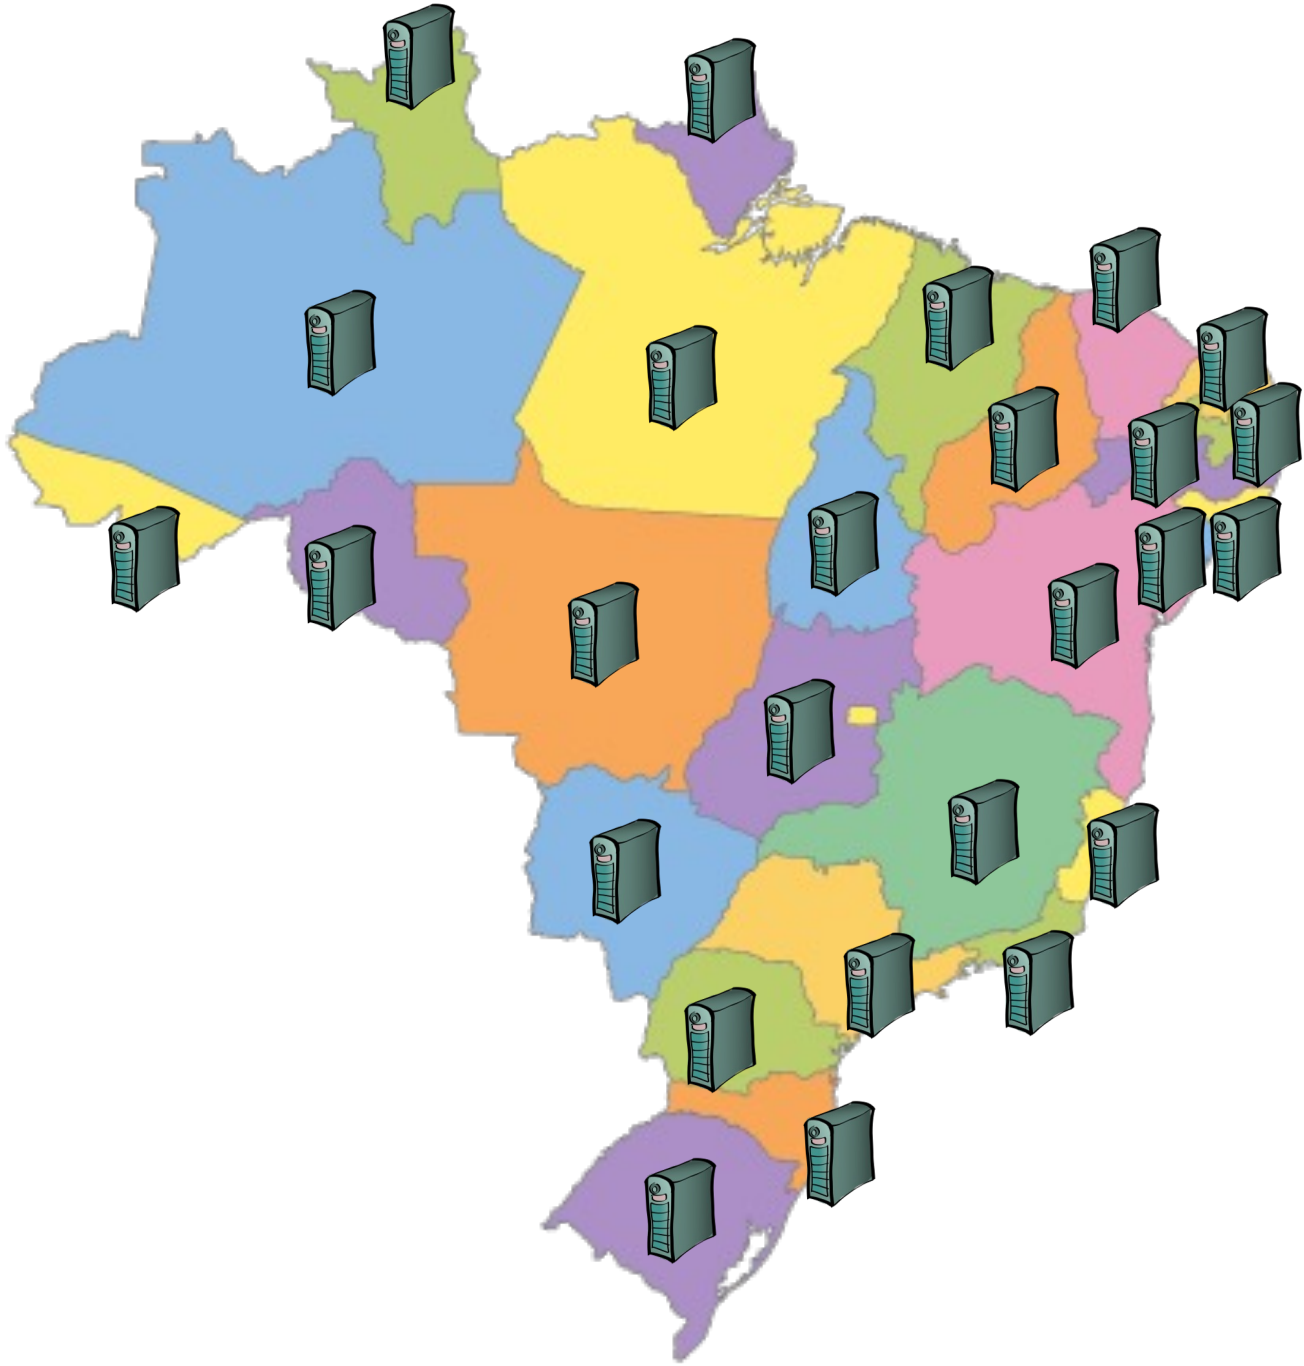

Monitoring a distributed system

- Periodically probe nodes
  - What frequency?
  - Liveliness vs overhead
- How to probe a node
  - What to check?
  - Thoroughness vs complexity

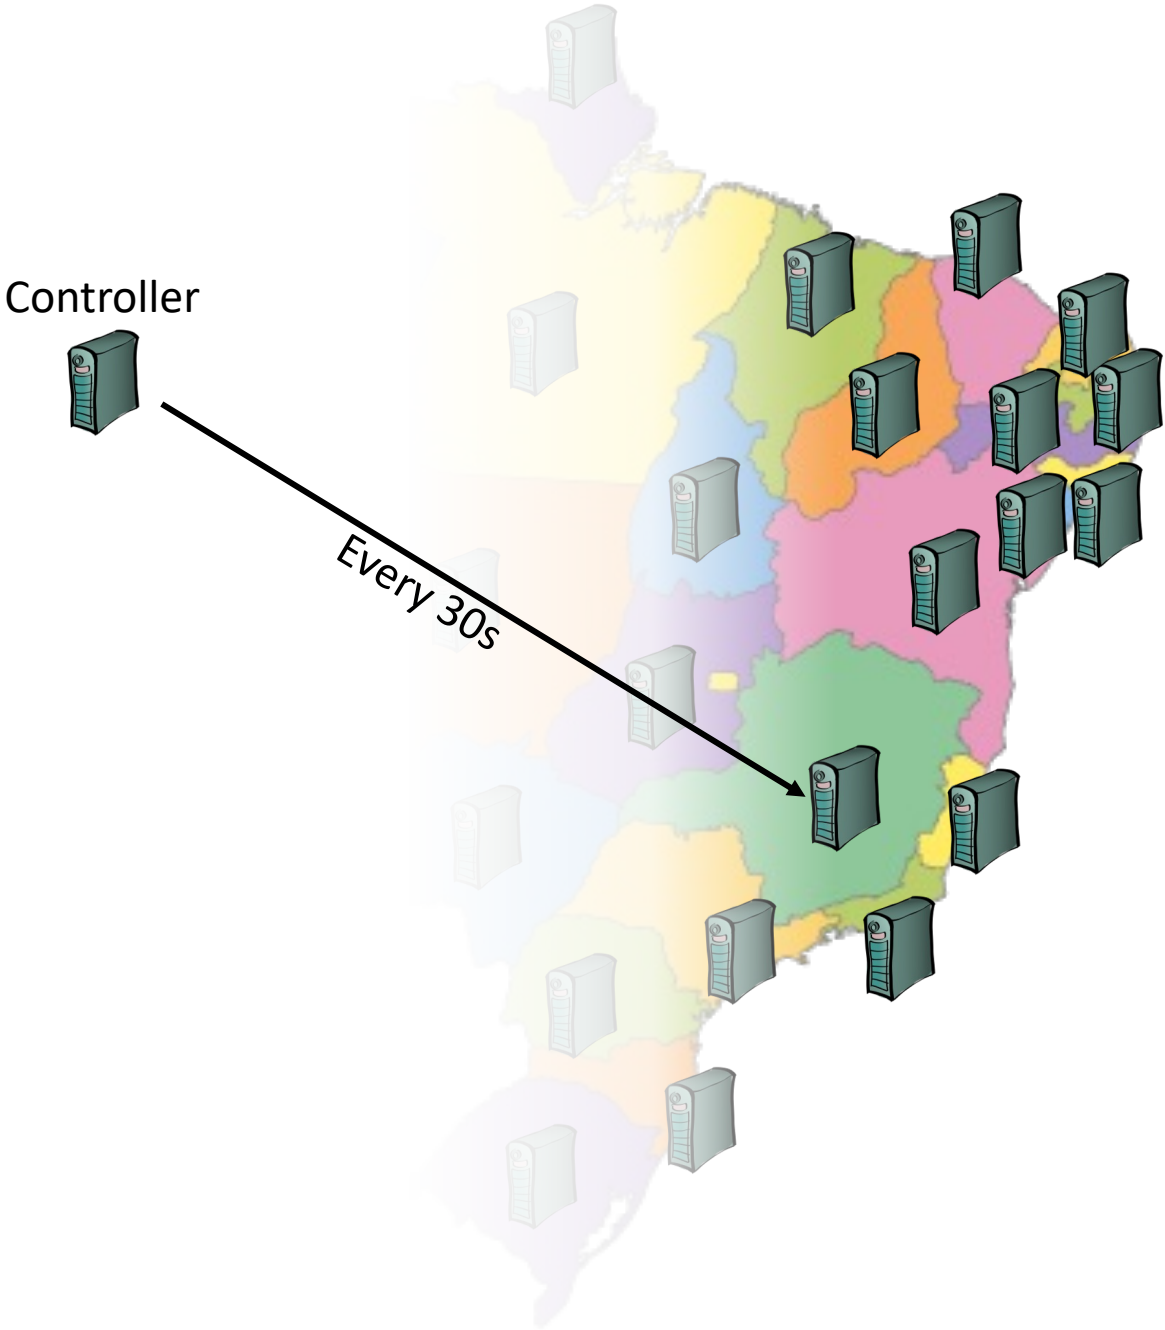

Monitoring a distributed system

- Periodically probe nodes
  - What frequency?
  - Liveliness vs overhead
- How to probe a node
  - What to check?
  - Thoroughness vs complexity
  - Ping → Check that the network is working
  - Simple request → Check network and app
  - Complex request → Network, app, and database

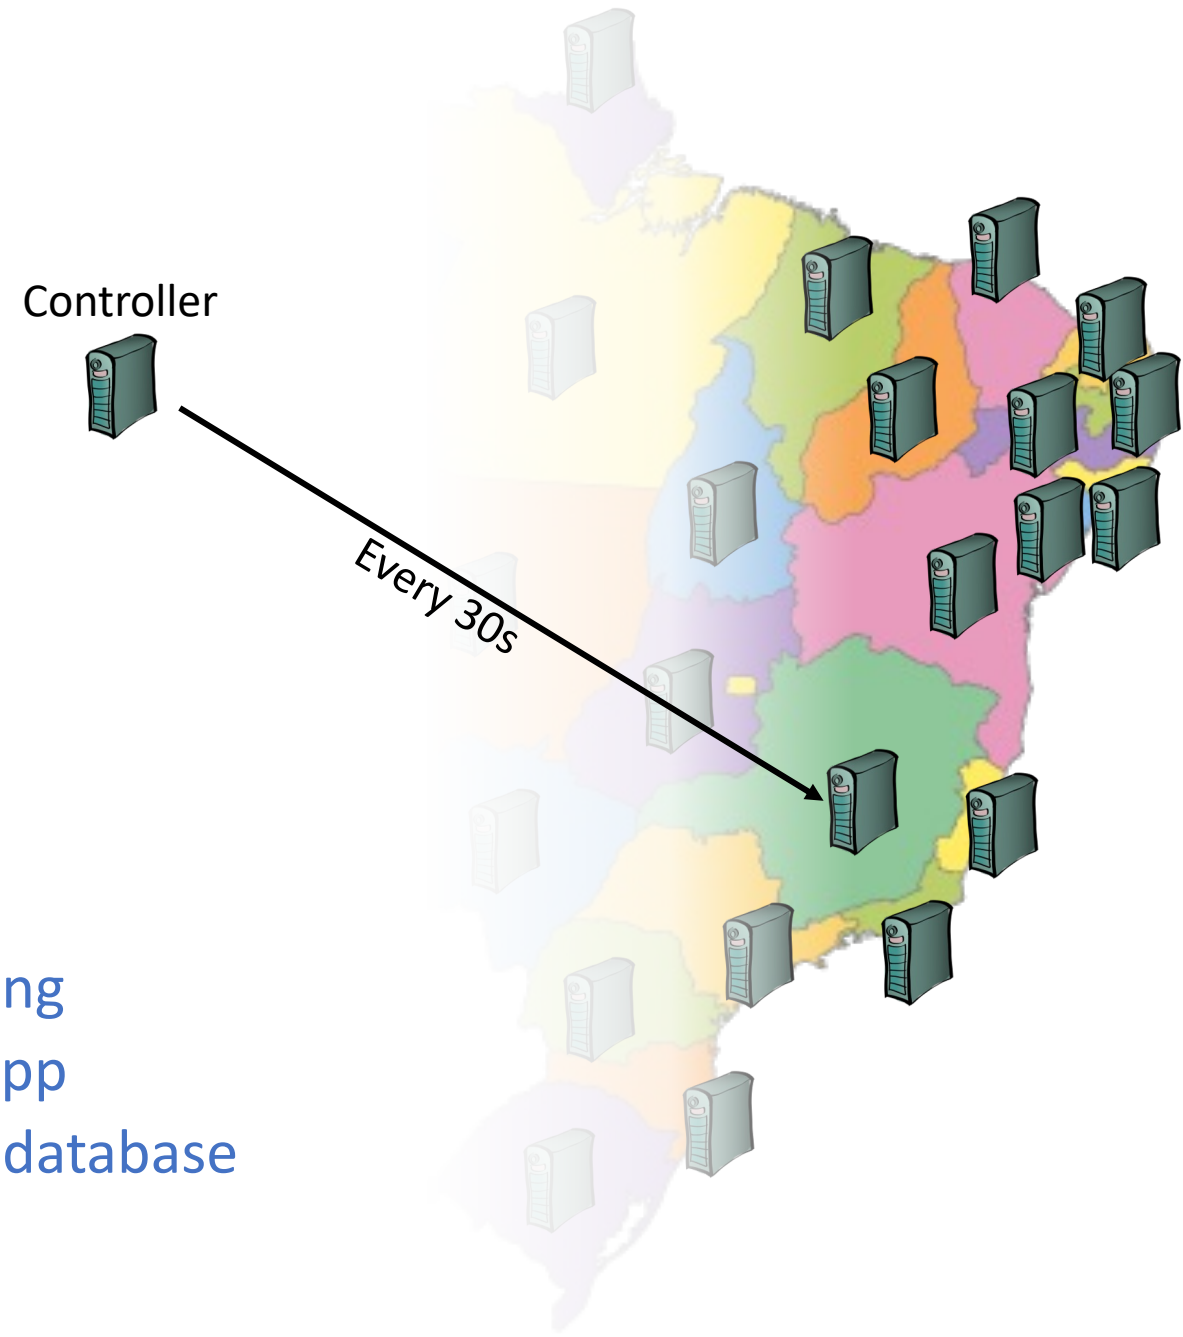

Monitoring a distributed system

- Periodically probe nodes
  - What frequency?
  - Liveliness vs overhead
- How to probe a node
  - What to check?
  - Thoroughness vs complexity
  - Ping → Check that the network is working
  - Simple request → Check network and app
  - Complex request → Network, app, and database

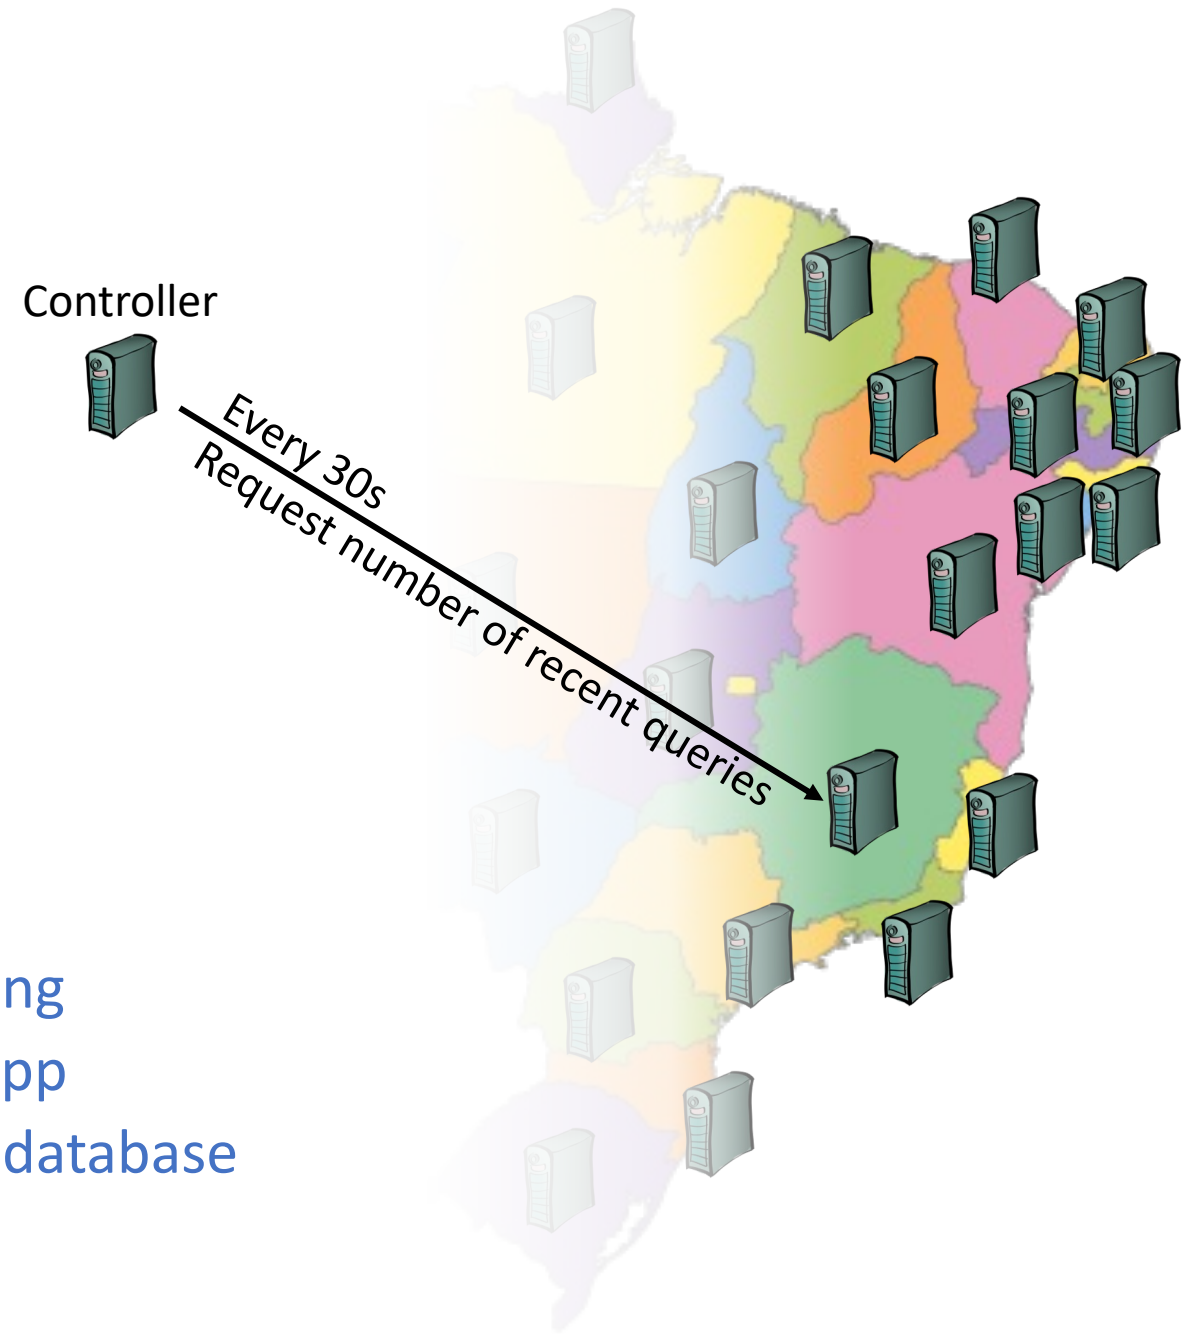

Monitoring a distributed system

- Periodically probe nodes
  - What frequency?
  - Liveliness vs overhead
- How to probe a node
  - What to check?
  - Thoroughness vs complexity
- Notify staff if failure detected

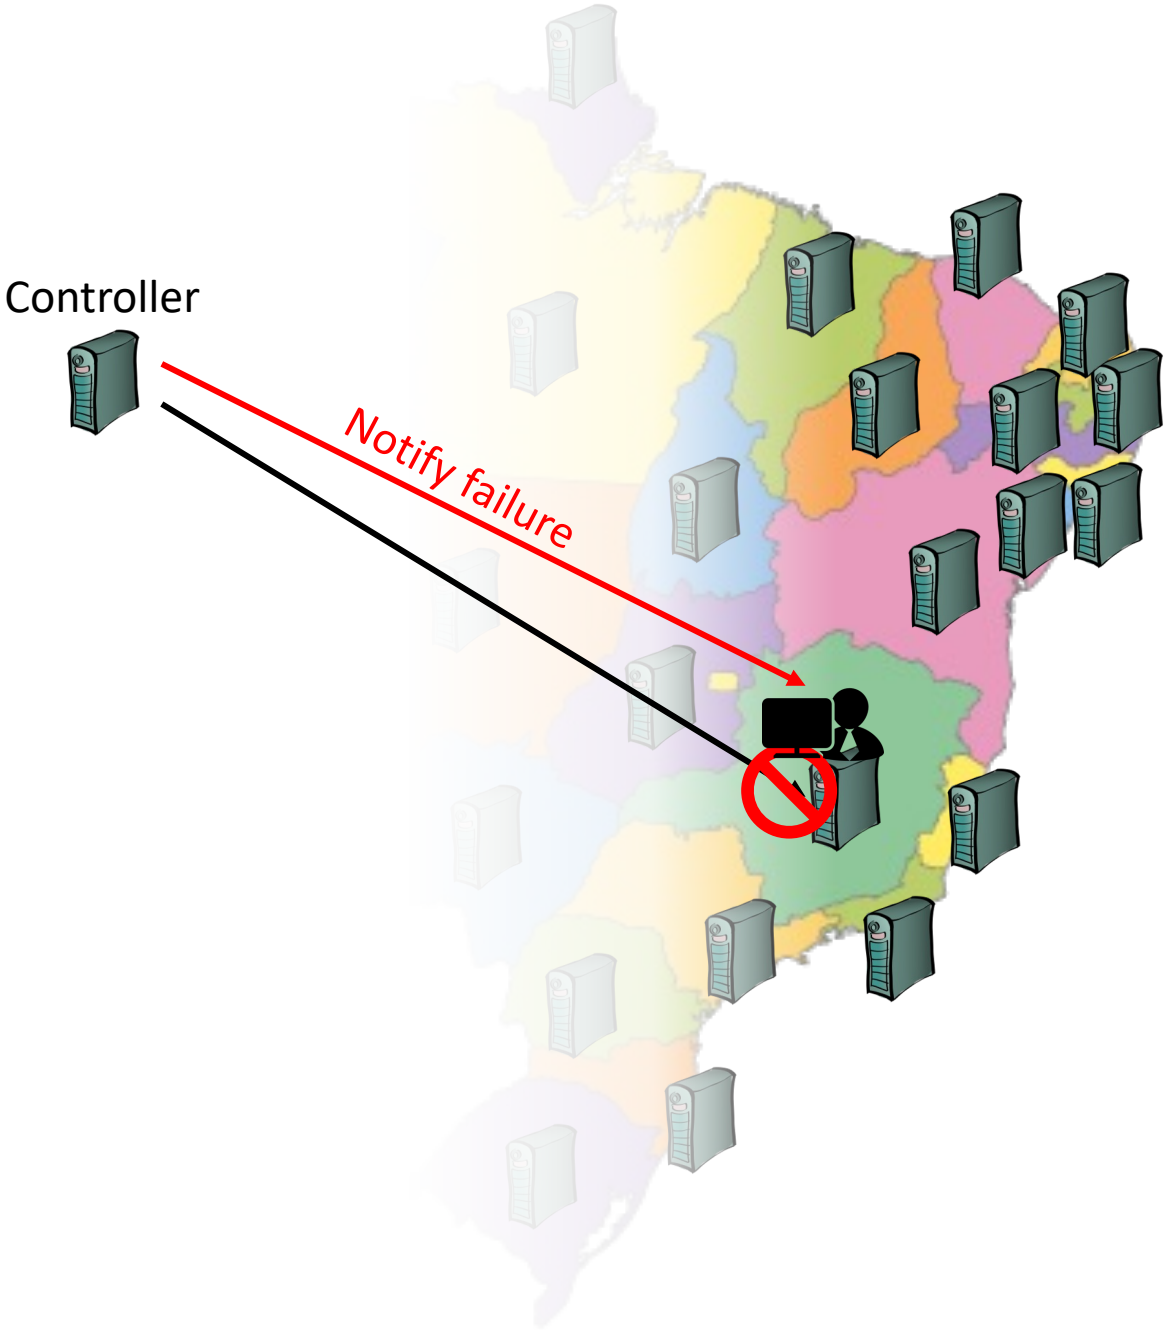

CompSci 401: Cloud Computing

Managing Cloud Applications

Prof. Ítalo Cunha

昆山杜克大学
DUKE KUNSHAN UNIVERSITY

Differences between distributed systems and cloud-native applications

Distributed System

- Monolithic applications on each node
- Multiple instances of one application

Cloud Native

- Containers/pods in each microservice
- Multiple instances of multiple microservices

Controller needs to handle a significantly larger number of instances.

Differences between distributed systems and cloud-native applications

Distributed System

- Monolithic applications on each node
- Multiple instances of one application
- Static set of instances
- Persistent applications

Cloud Native

- Containers/pods in each microservice
- Multiple instances of multiple microservices
- Instances may move (placement)
- Ephemeral containers

Cannot know which instances to monitor in advance.

Differences between distributed systems and cloud-native applications

Distributed System

- Monolithic applications on each node
- Multiple instances of one application
- Static set of instances
- Persistent applications
- Fixed number of instances

Cloud Native

- Containers/pods in each microservice
- Multiple instances of multiple microservices
- Instances may move (placement)
- Ephemeral containers
- Autoscaling creates/destroys instances

New instances may appear, and instances may disappear. Disappearing instances ***is OK!***

CompSci 401: Cloud Computing

Control Loops

Prof. Ítalo Cunha

昆山杜克大学
DUKE KUNSHAN UNIVERSITY

Control loop

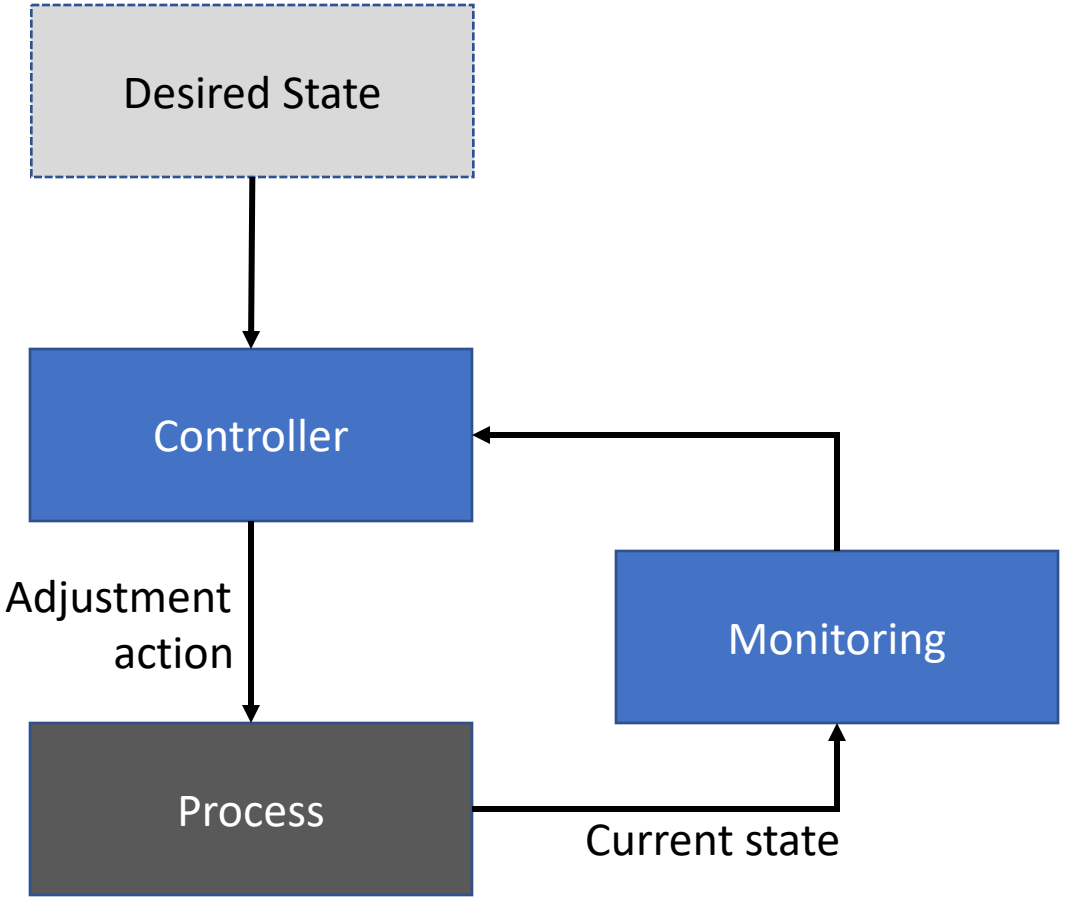

Control loop

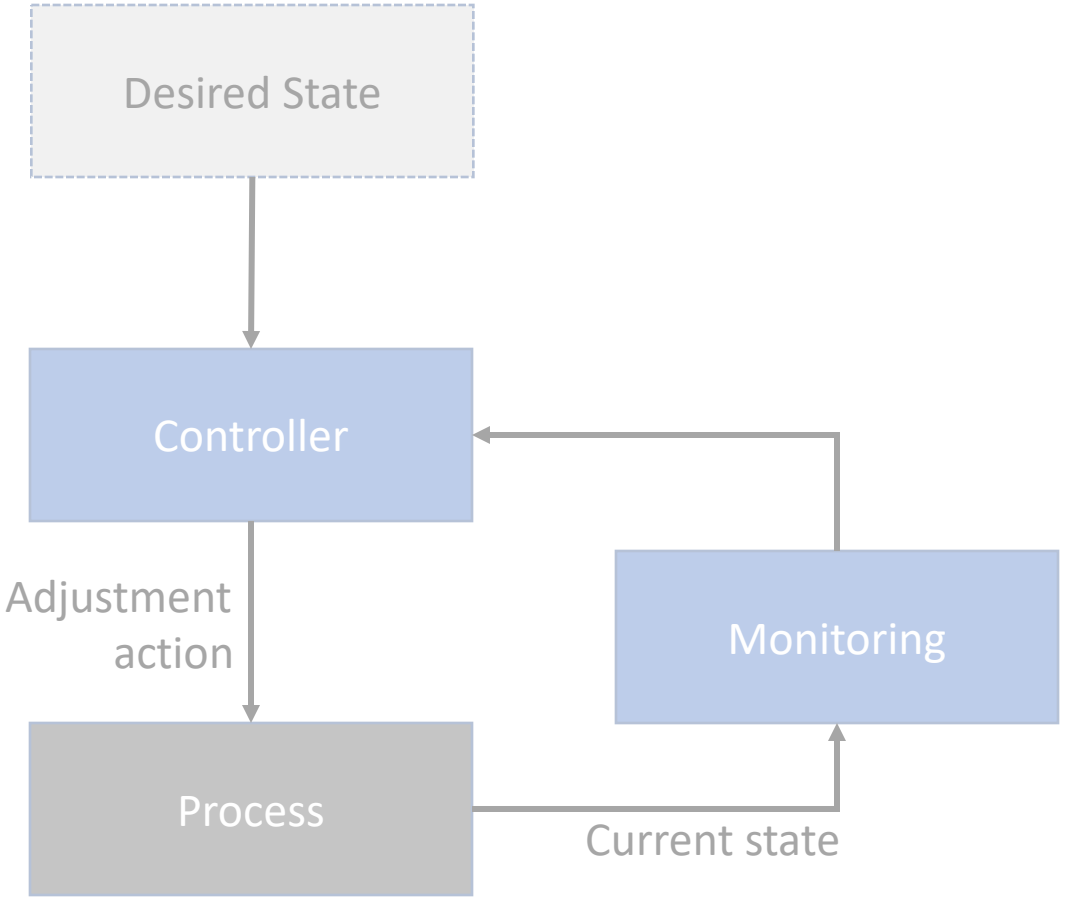

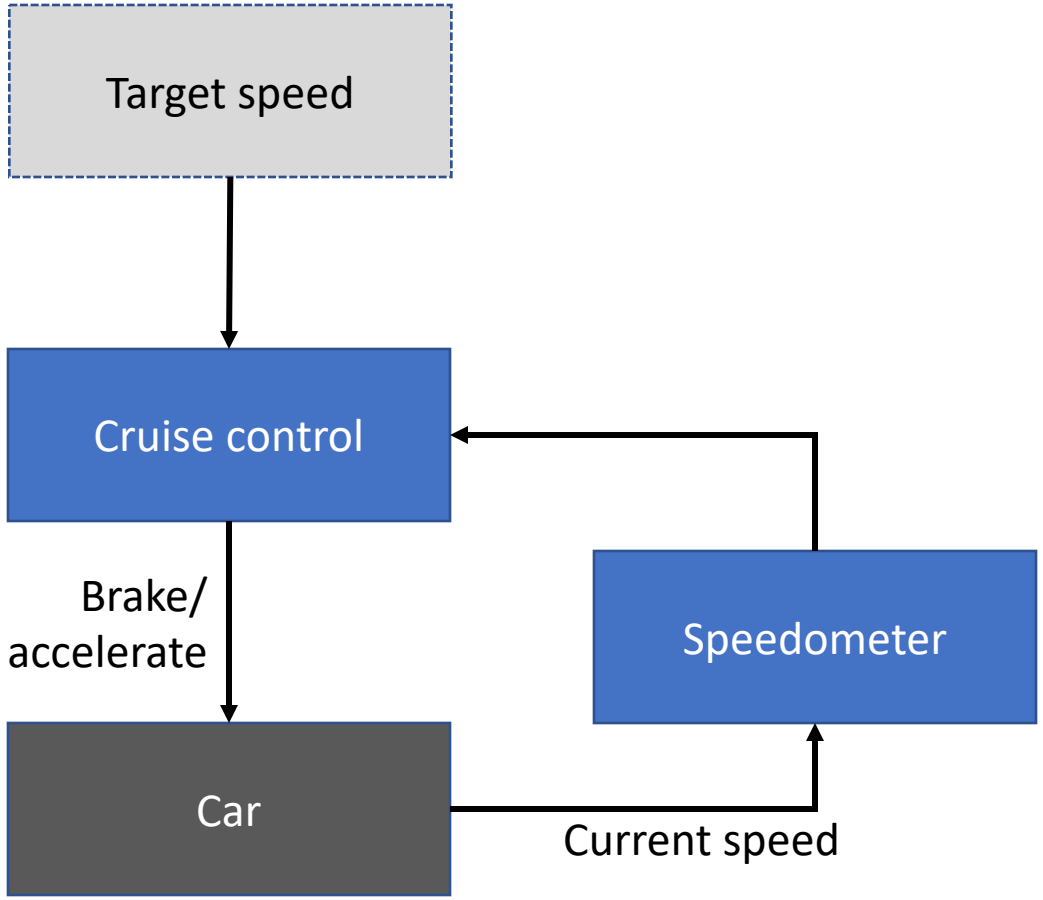

Control loop

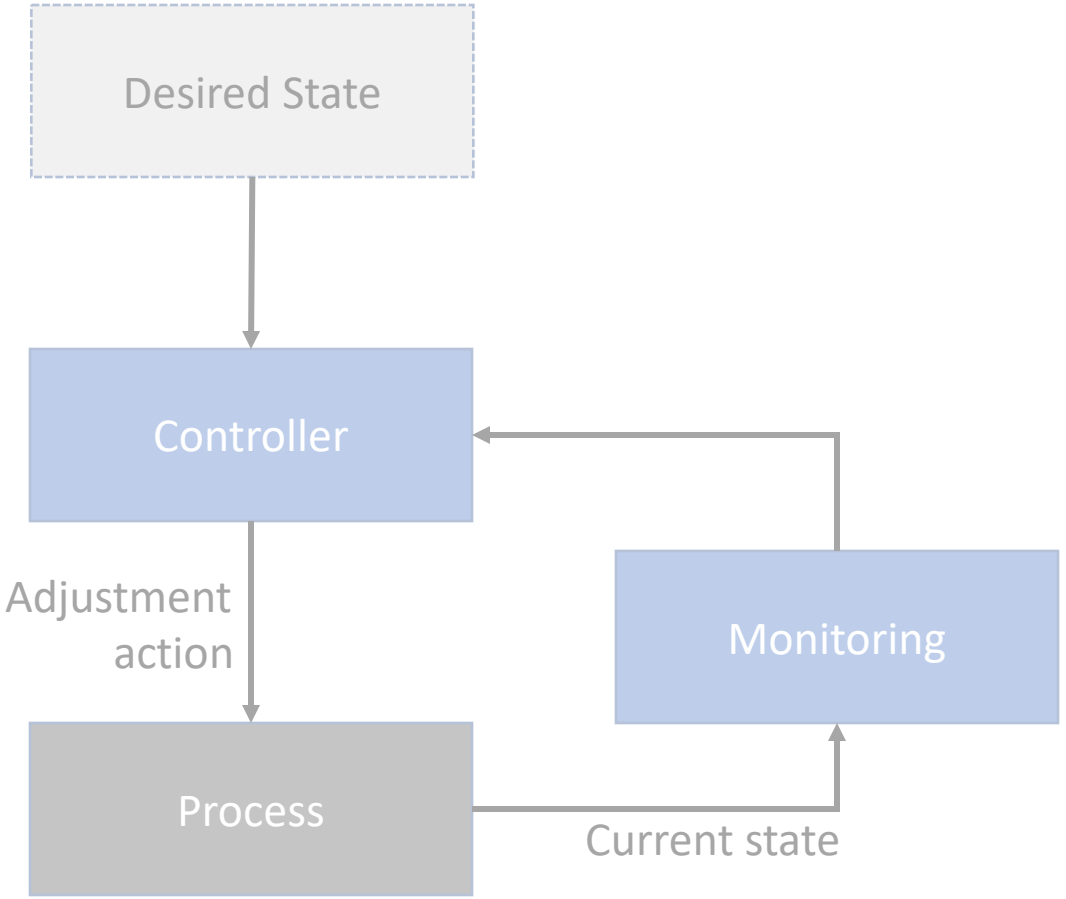

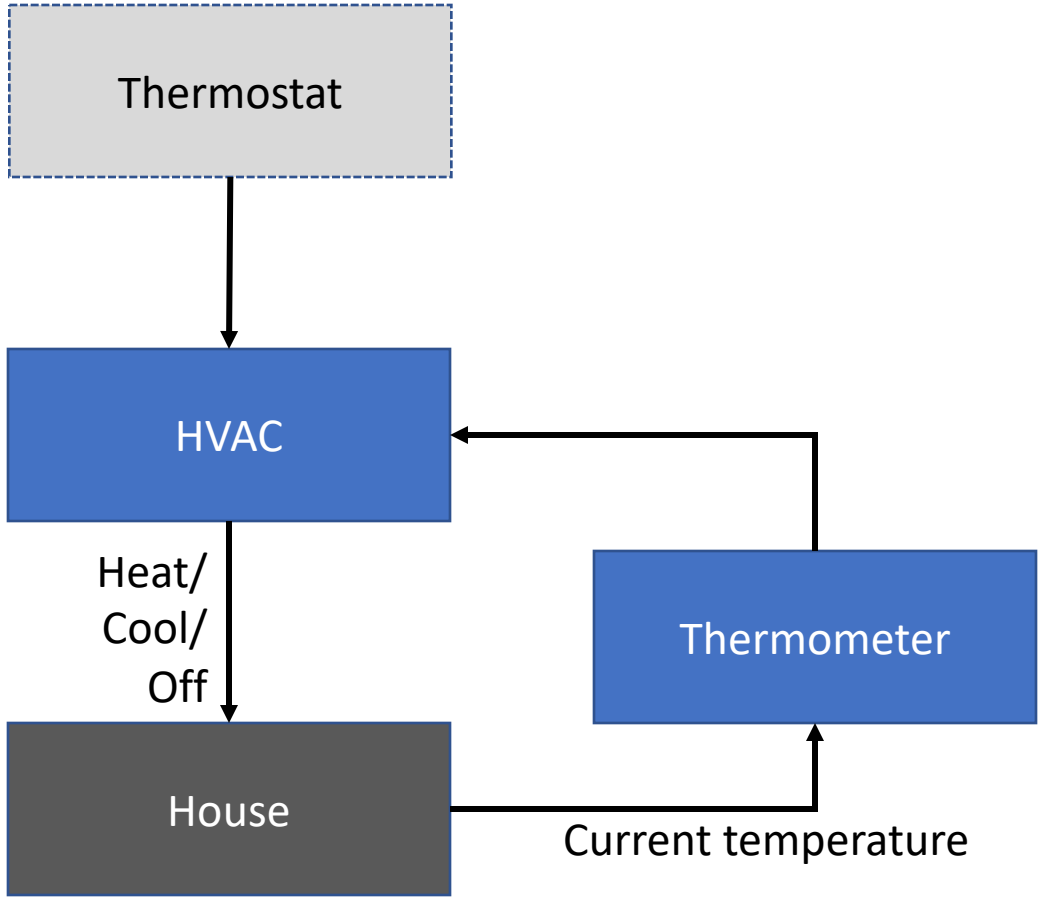

Control loop

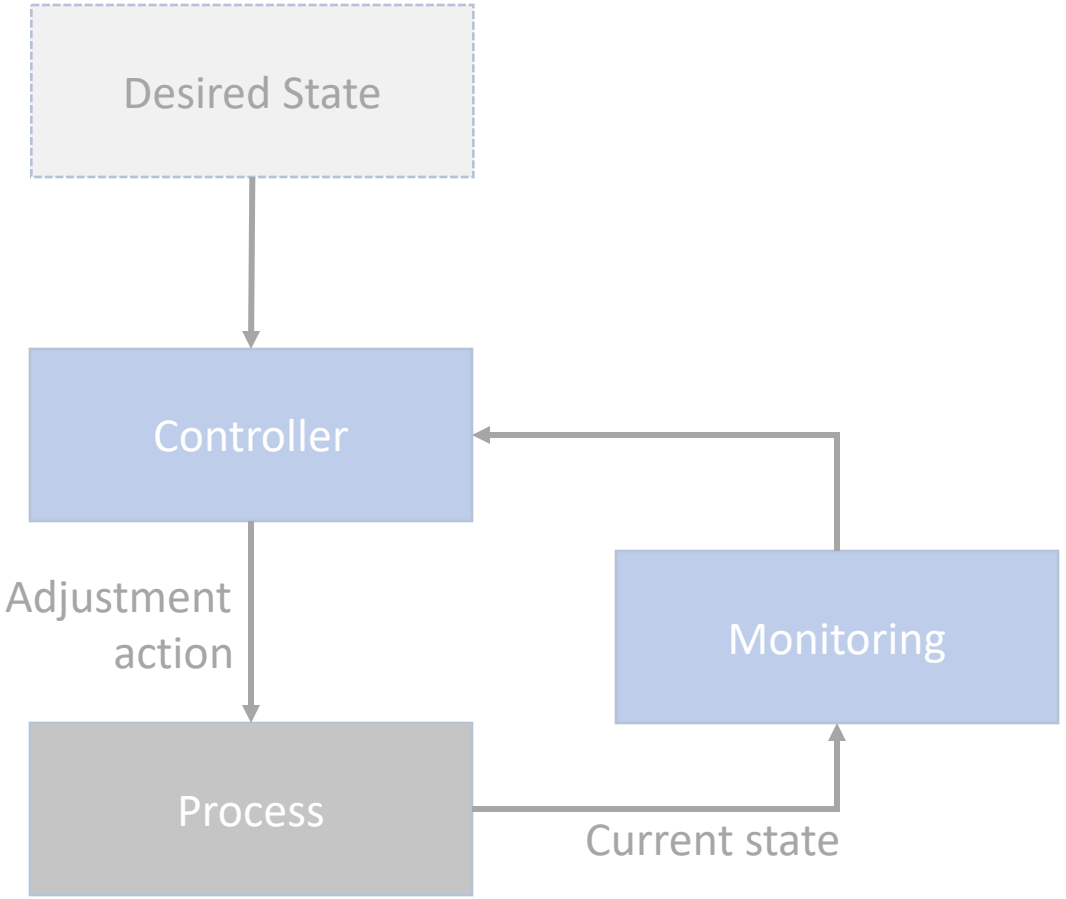

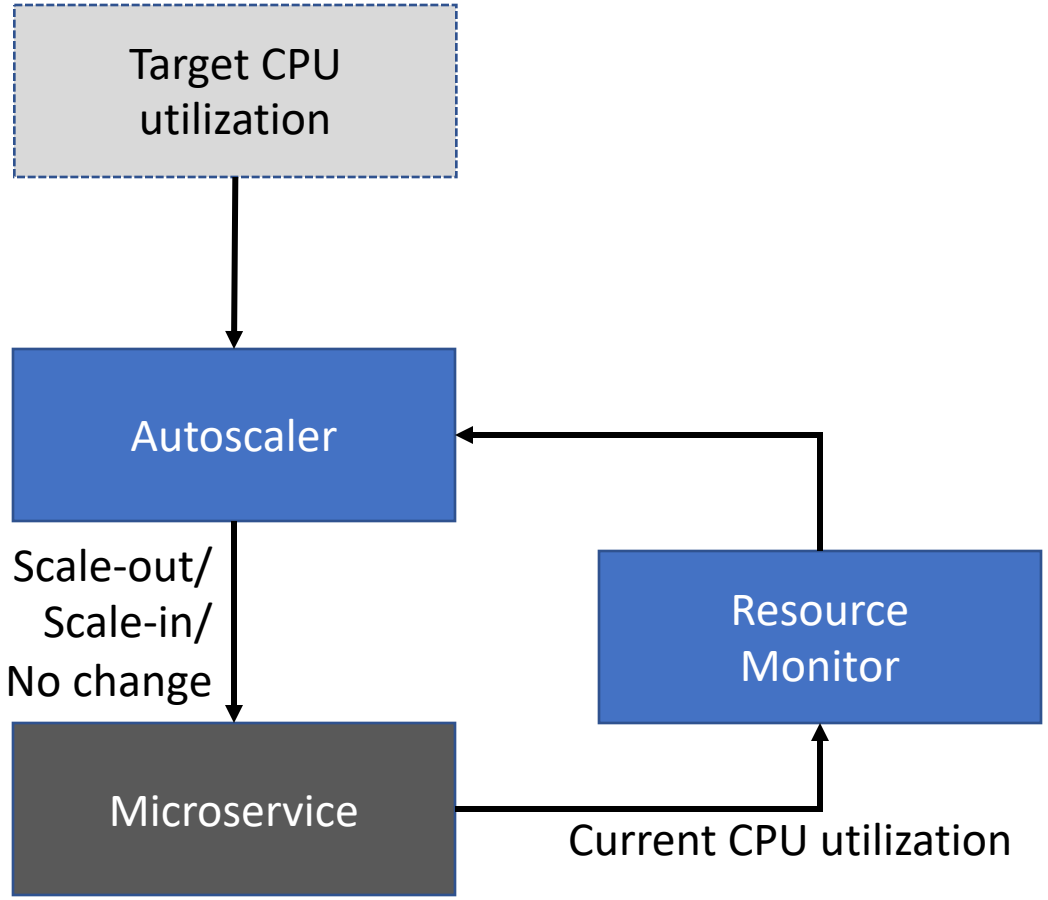

Declarative interface

- We say what is the desired state
- *Not how to achieve it*

Declarative interface

- We say what is the desired state
- *Not how to achieve it*

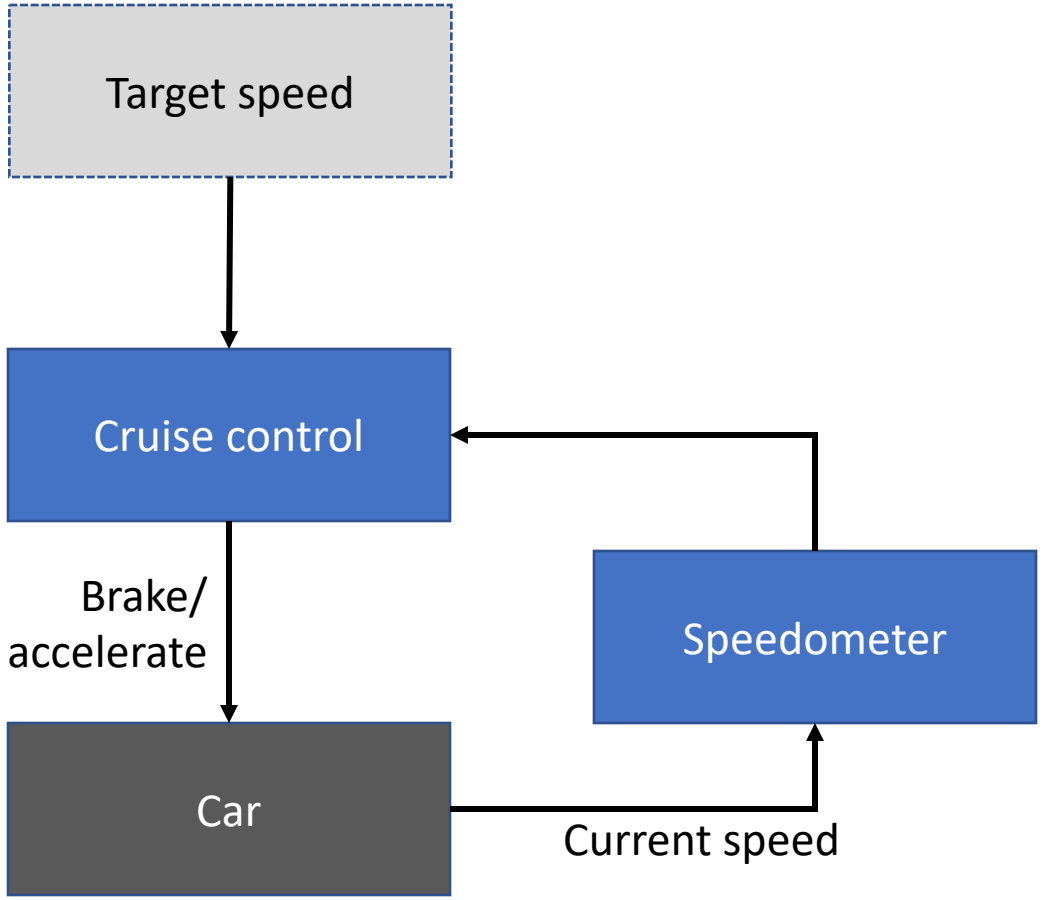

Control loop delay

- How frequently to measure state and perform control actions?
  - Application dependent

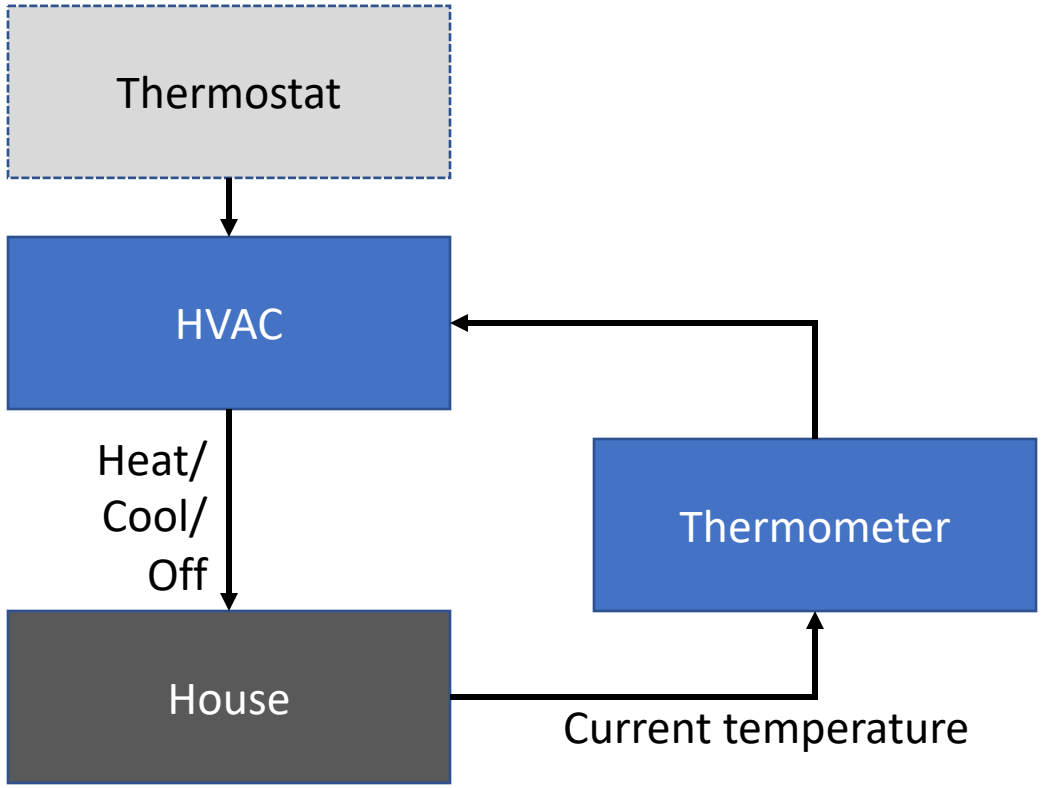

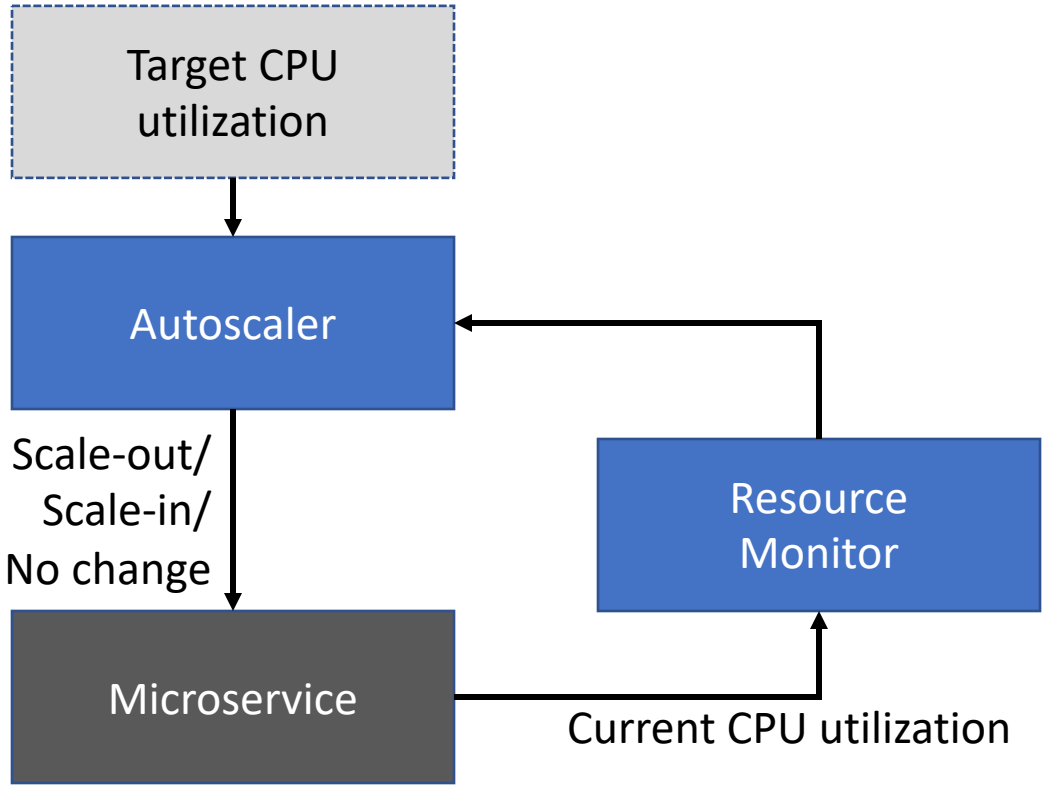

Control loop delay

- How frequently to measure state and perform control actions?
  - Application dependent

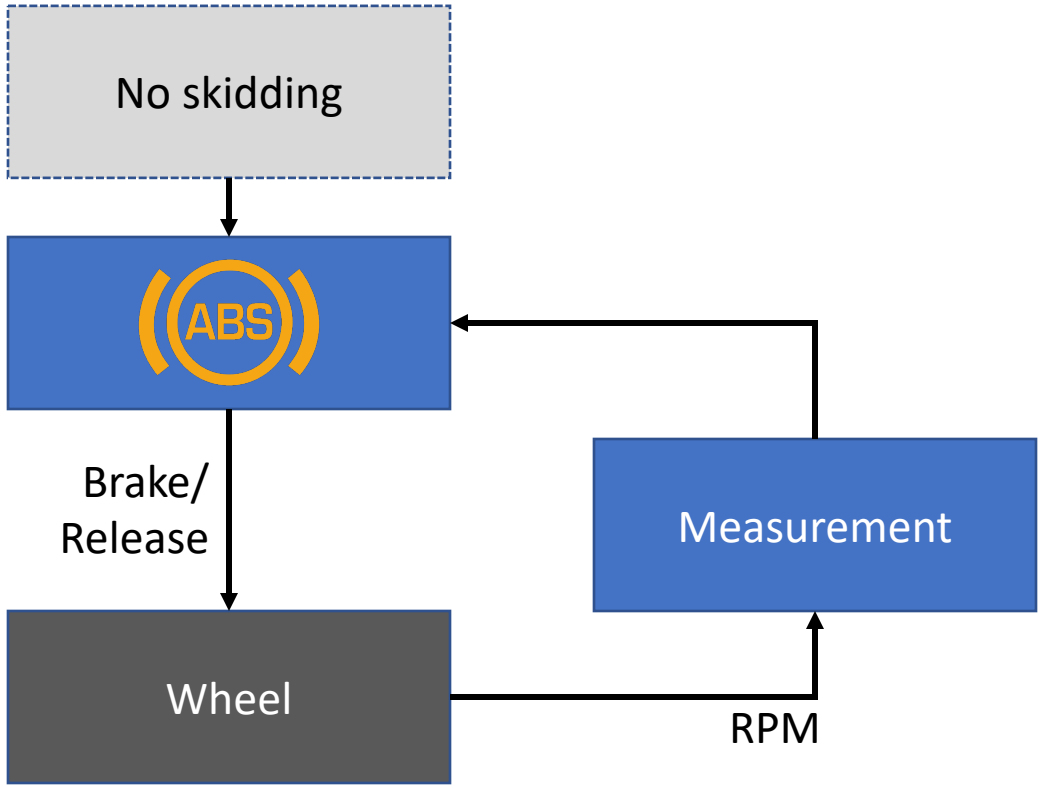

Hysteresis

- State may take time to change, control system needs to consider

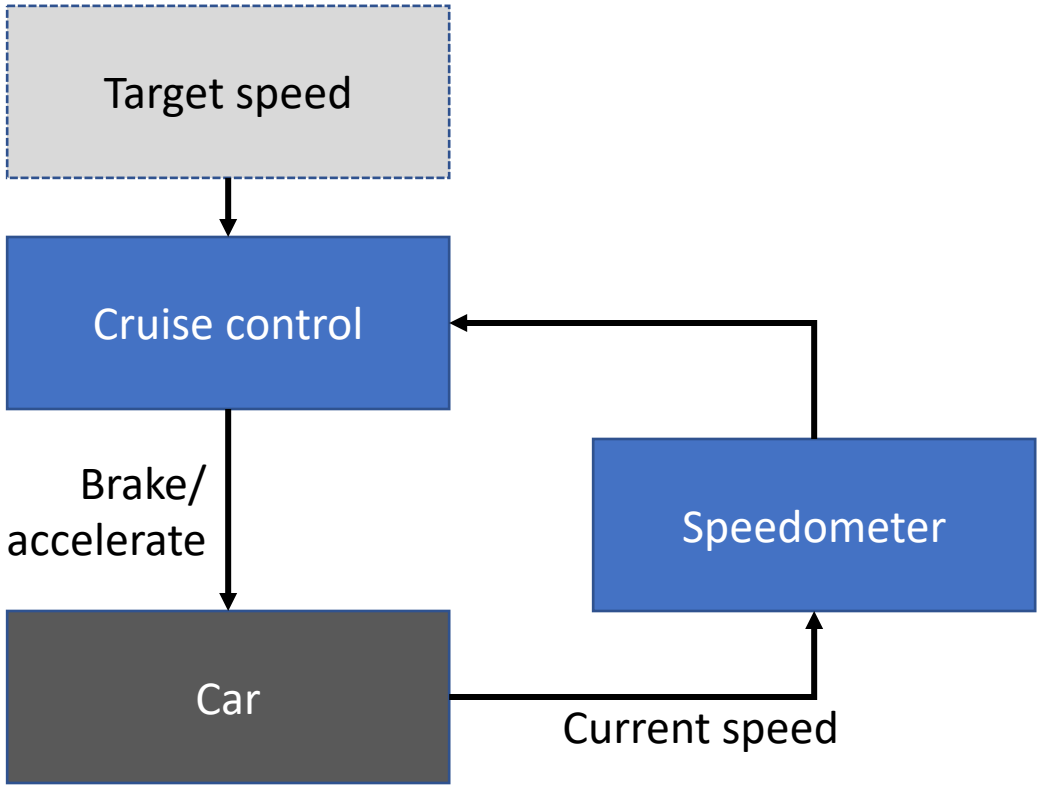

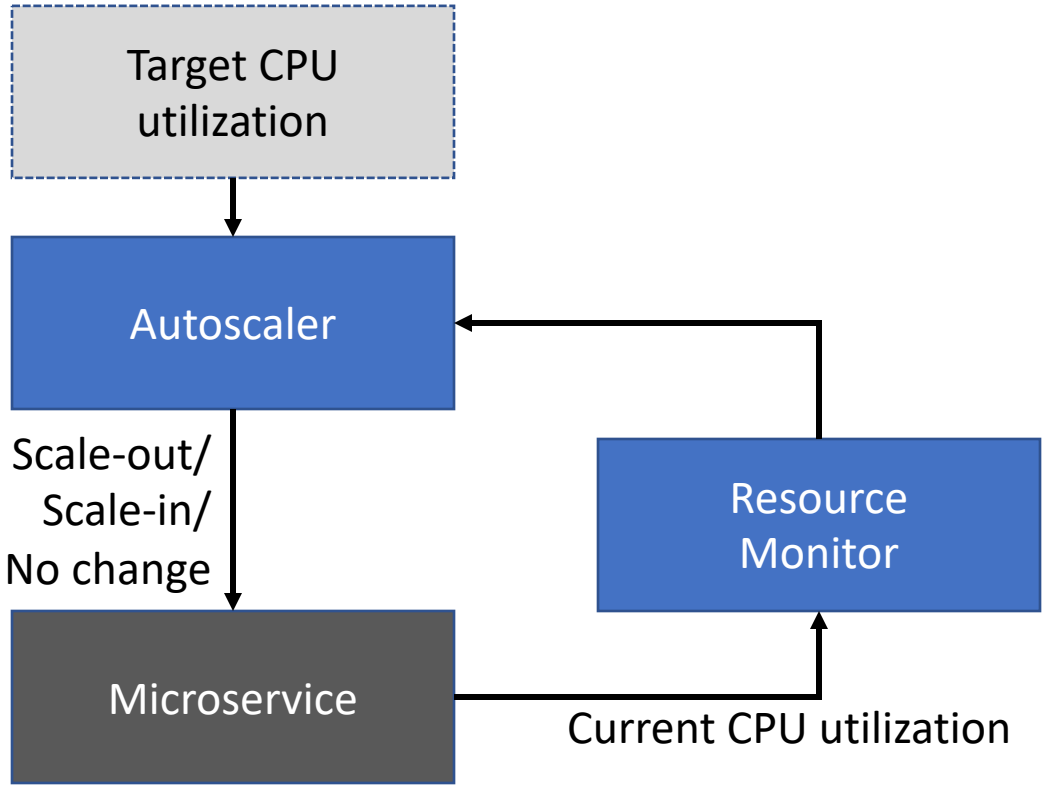

Hysteresis may lead to undesirable conditions

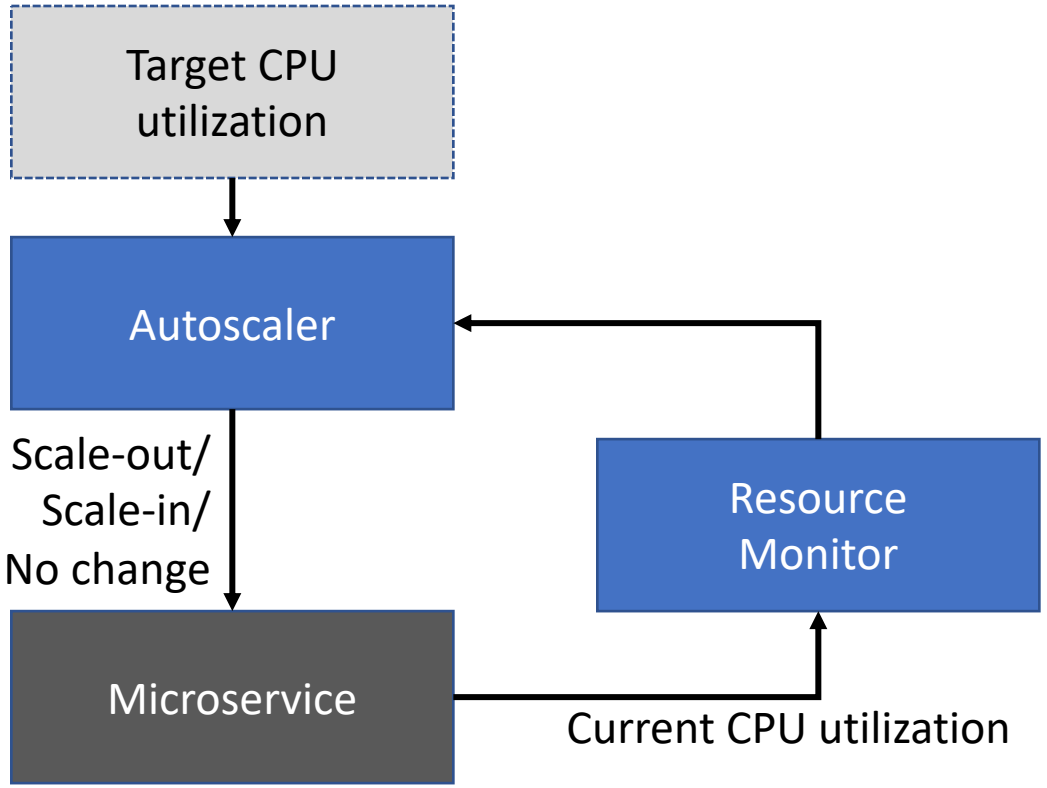

Sequence of events

- T1: Resource monitor tells autoscaler that the microservice is overloaded
- T2: Scale-out 2 instances
- T3: New instances start

Hysteresis may lead to undesirable conditions

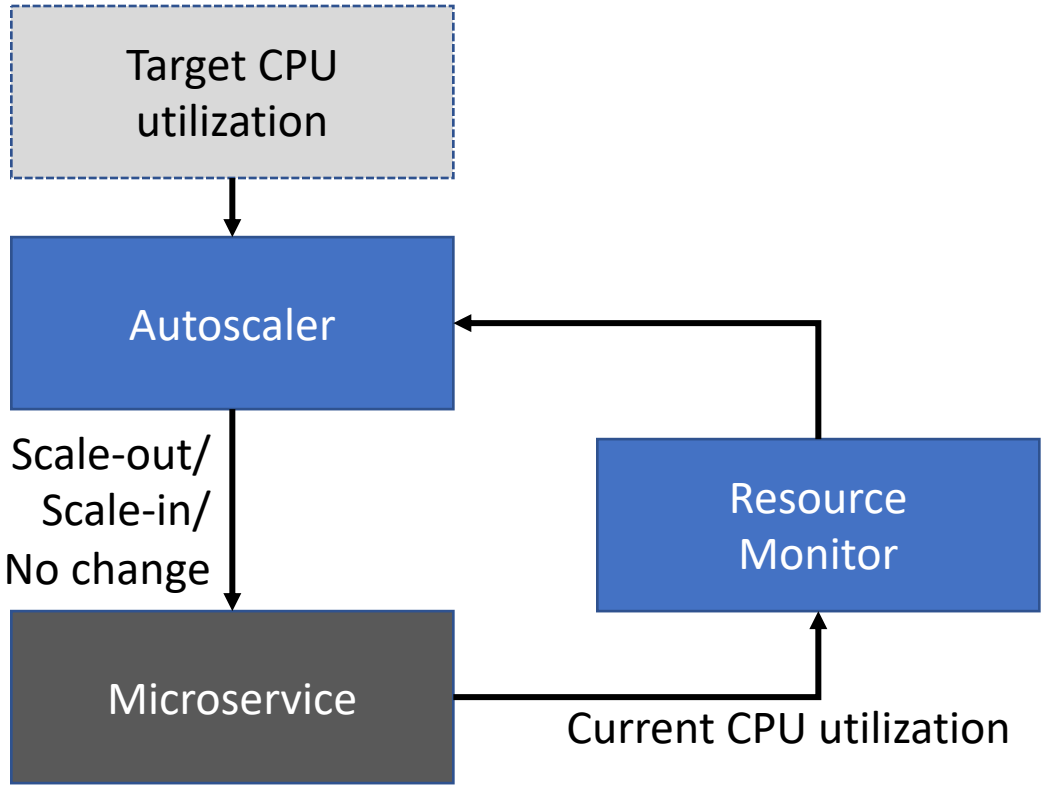

Sequence of events

- T1: Resource monitor tells autoscaler that the microservice is overloaded
- T2: Scale-out 2 instances
- T3: New instances start
- T4: Load balancer starts to distribute load
- T5: Resource monitor tells autoscaler that the microservice is overloaded
- T6: Scale-out 2 instances

Hysteresis may lead to undesirable conditions

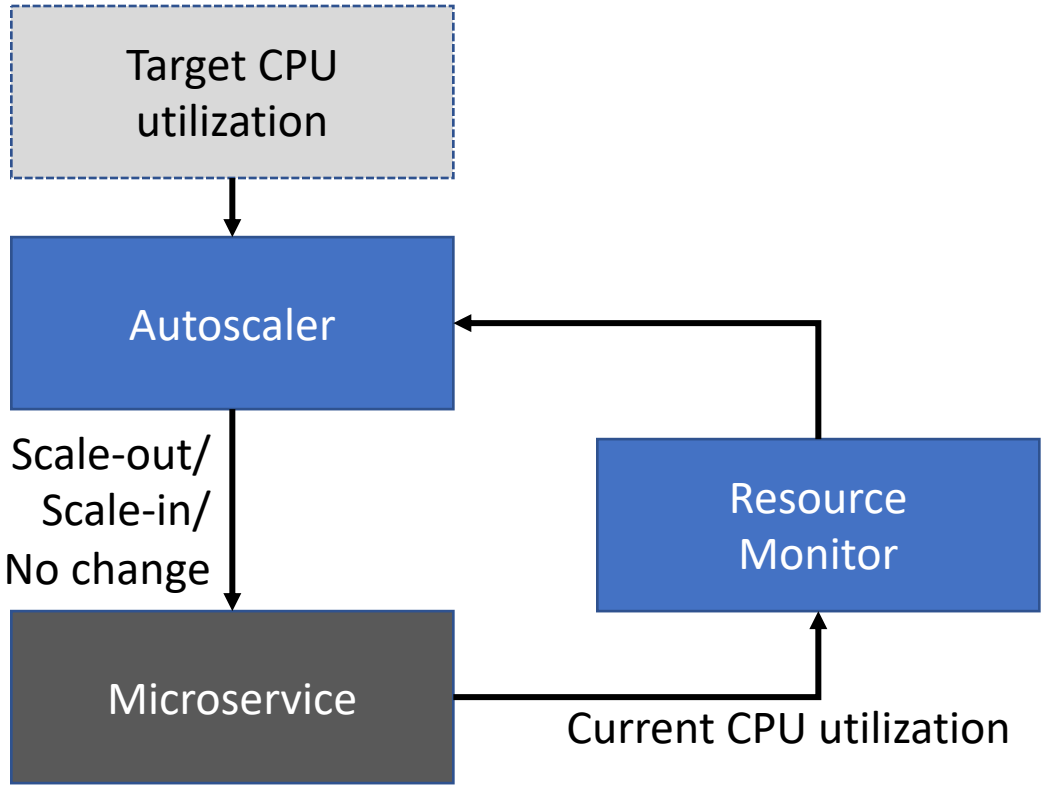

Sequence of events

- T1: Resource monitor tells autoscaler that the microservice is overloaded
- T2: Scale-out 2 instances
- T3: New instances start
- T4: Load balancer starts to distribute load
- T5: Resource monitor tells autoscaler that the microservice is overloaded
- T6: Scale-out 2 instances

+4 instances may be unecessary!

Hysteresis may lead to instability

Hysteresis may lead to instability

Hysteresis may lead to instability

Hysteresis may lead to instability

Hysteresis may lead to instability

Hysteresis may lead to instability

Hysteresis may lead to instability

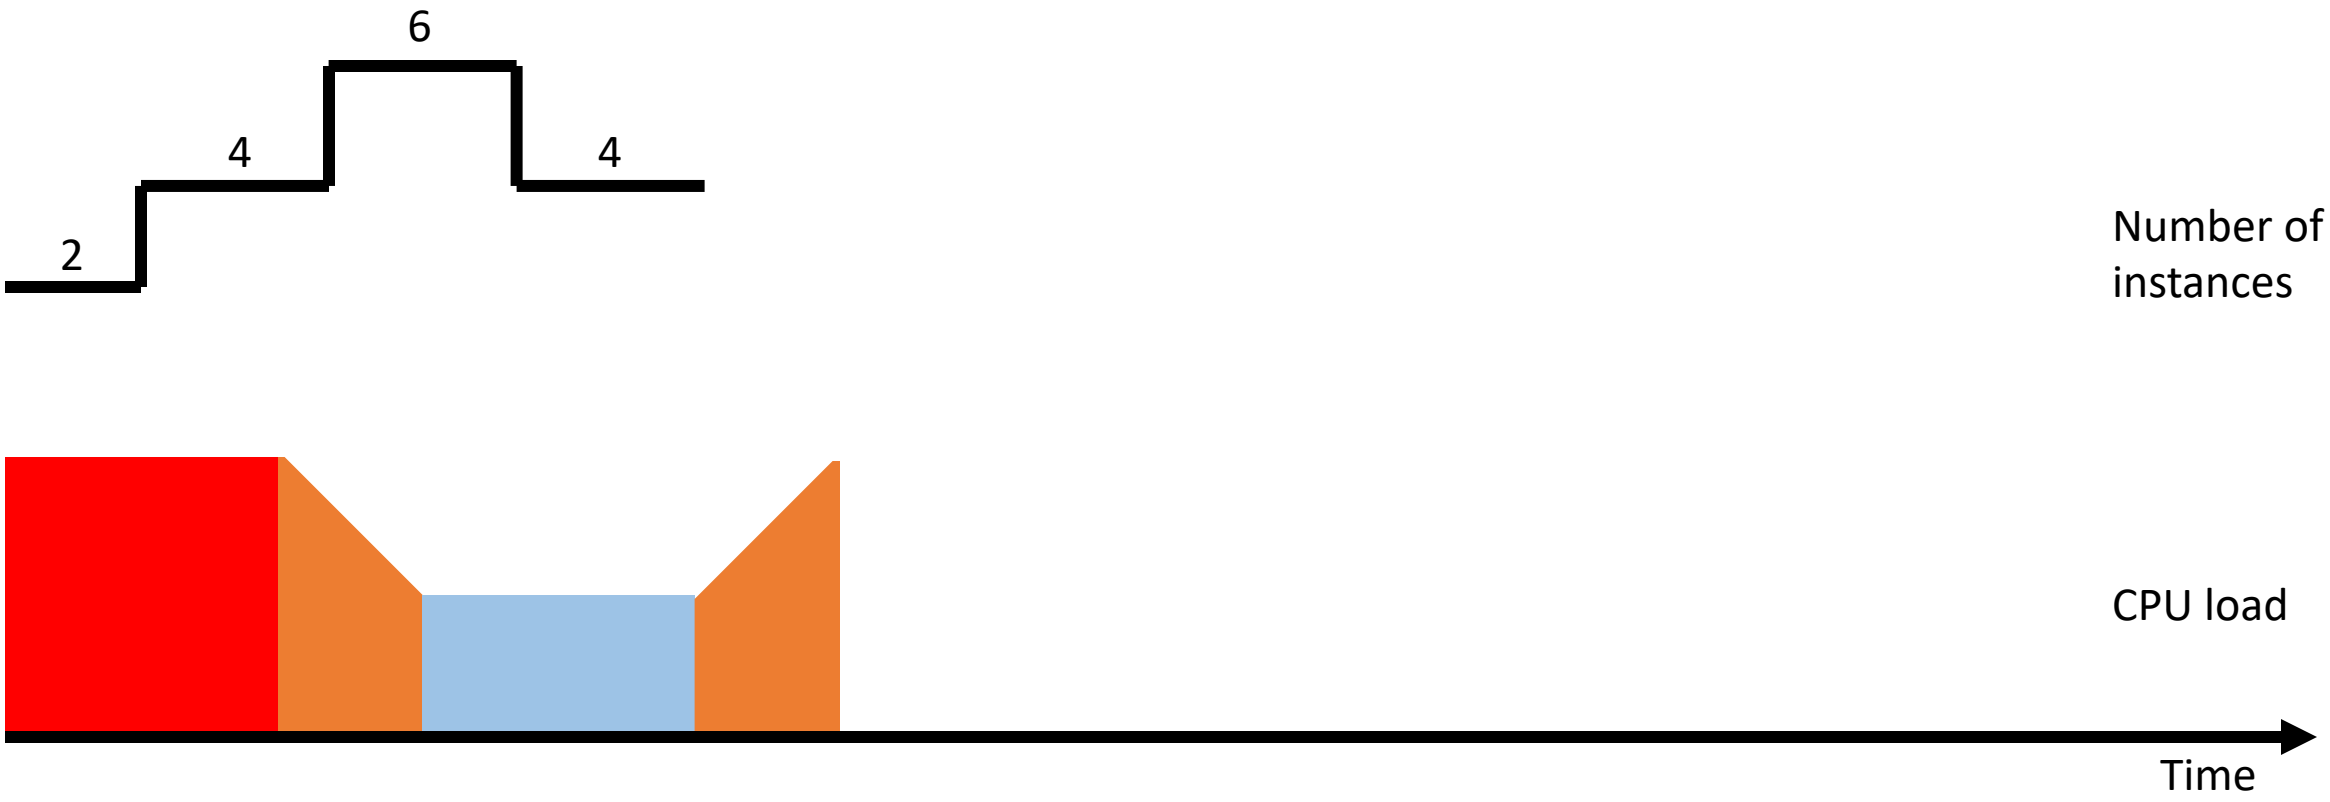

Hysteresis may lead to instability

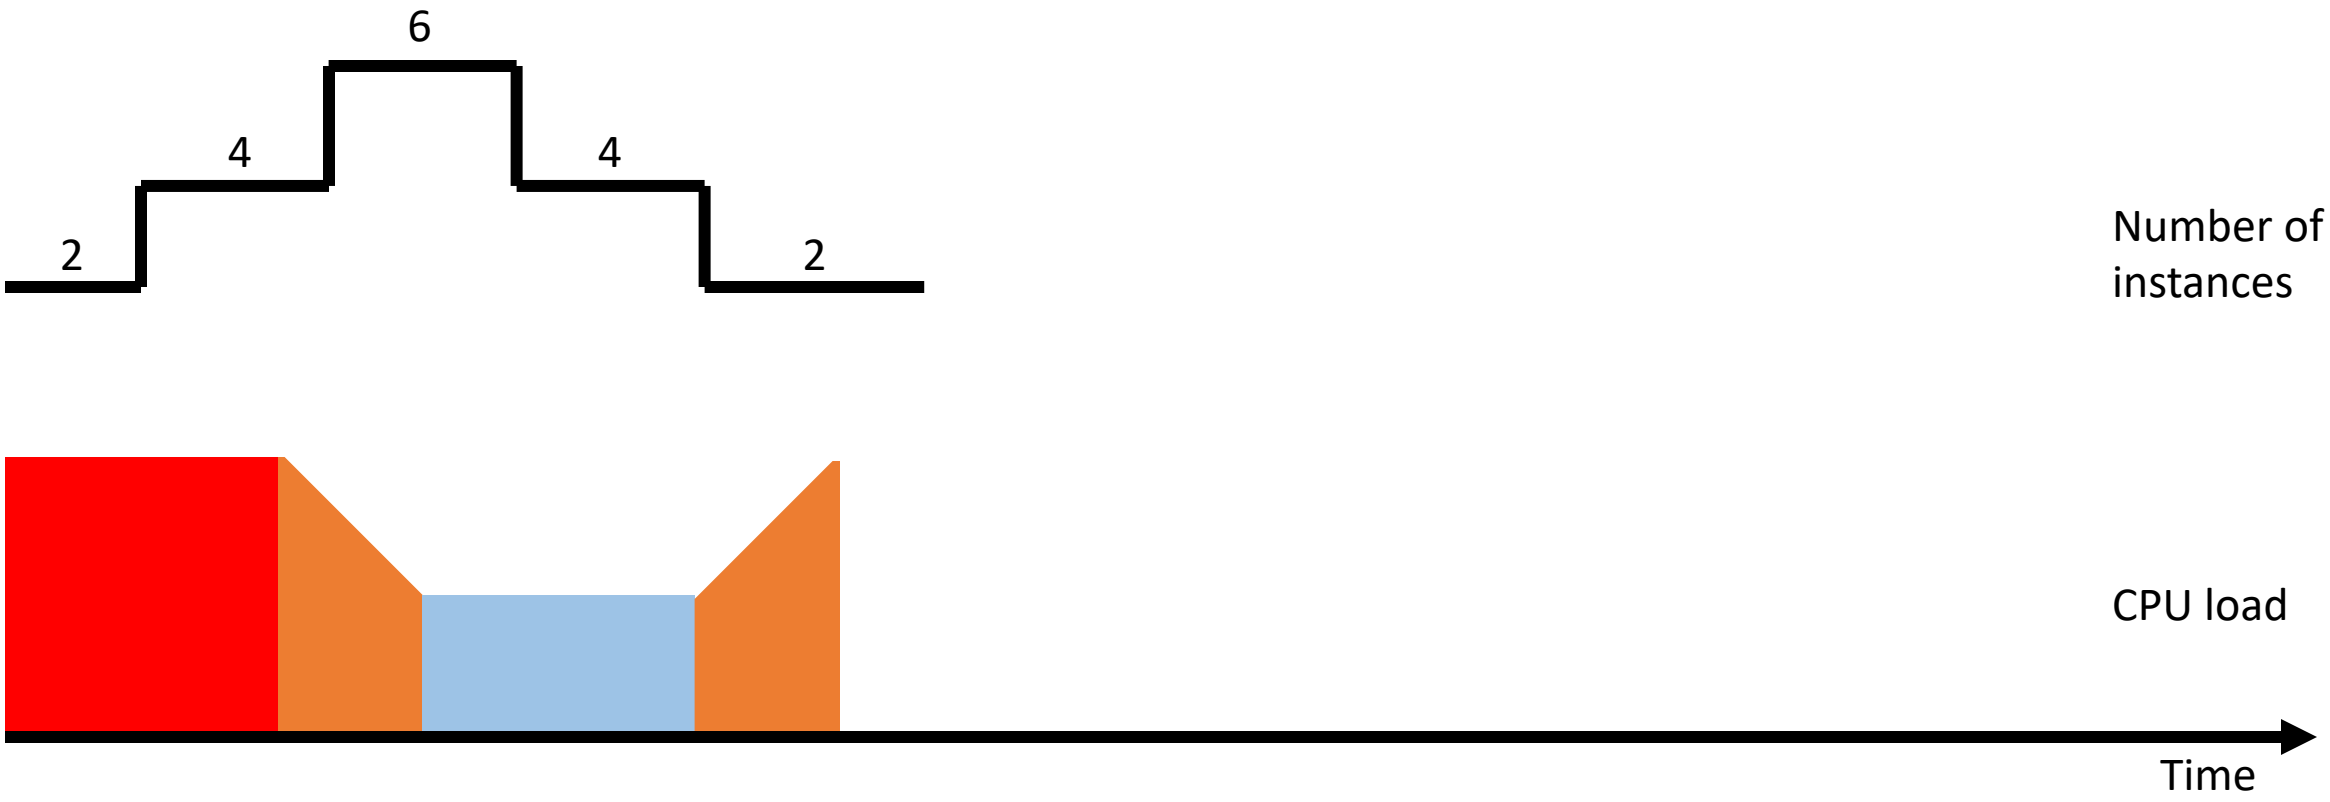

Hysteresis may lead to instability

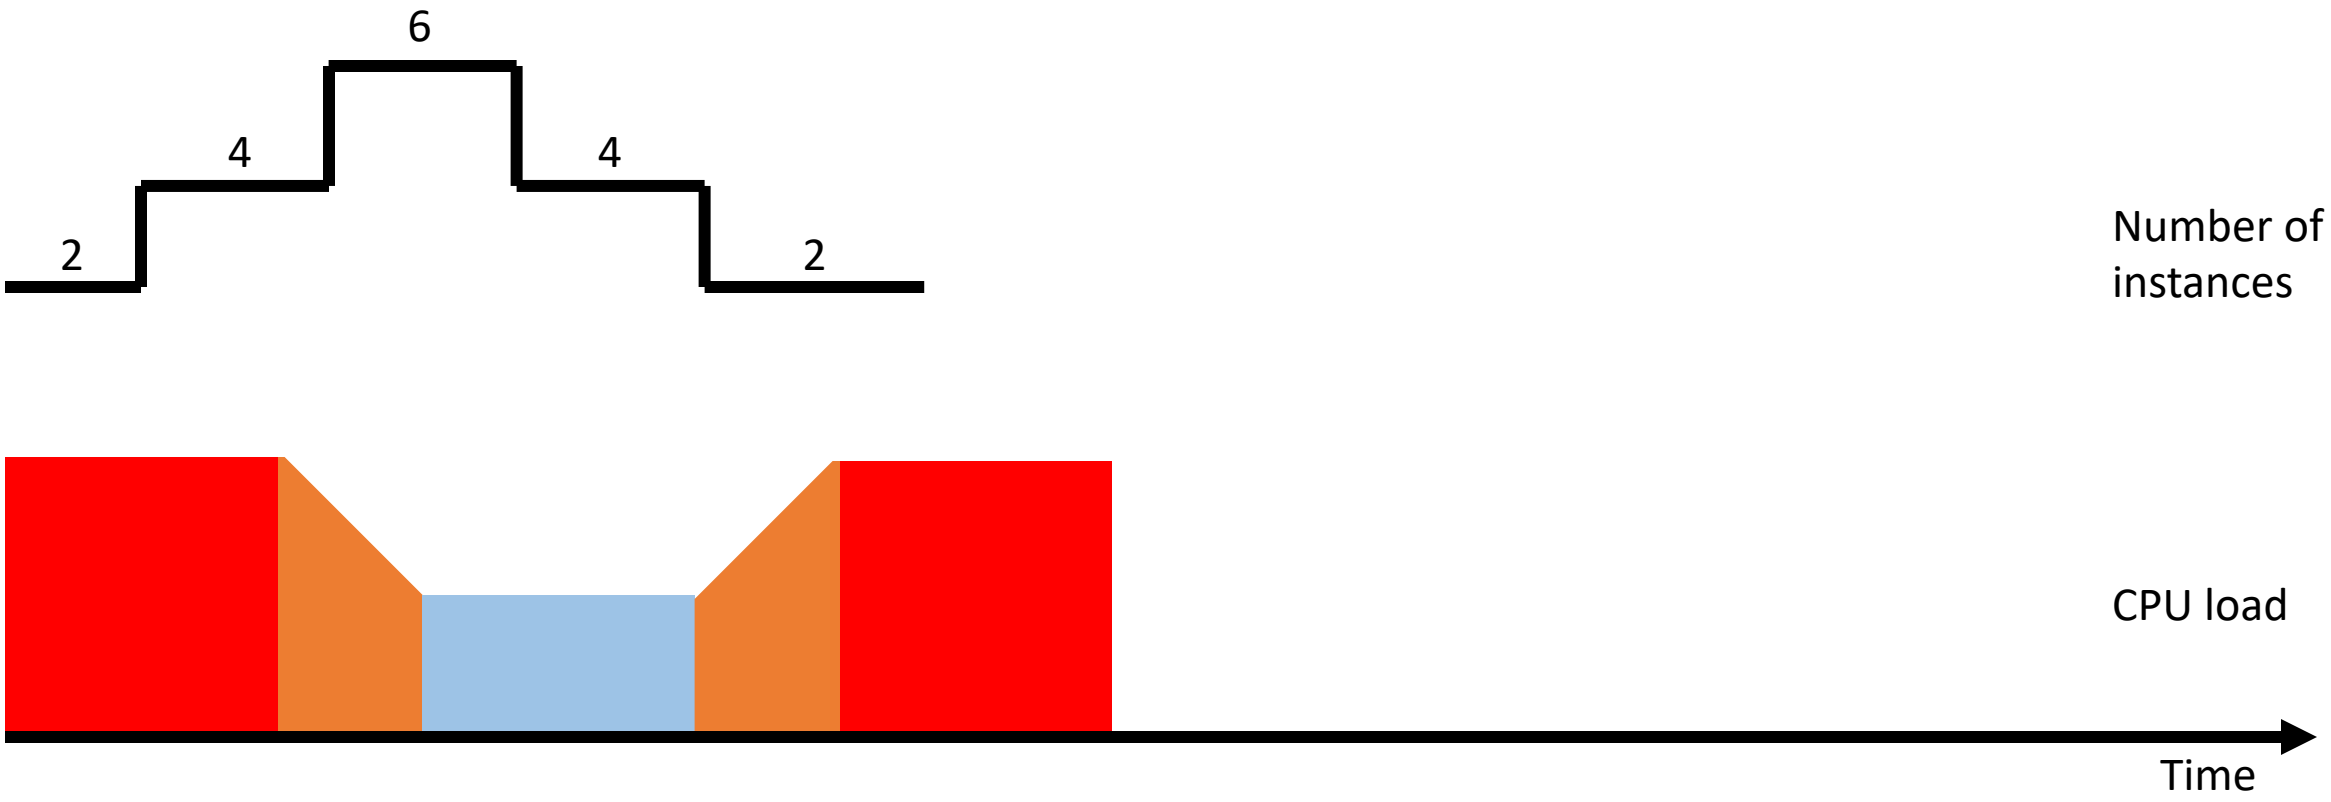

Hysteresis may lead to instability

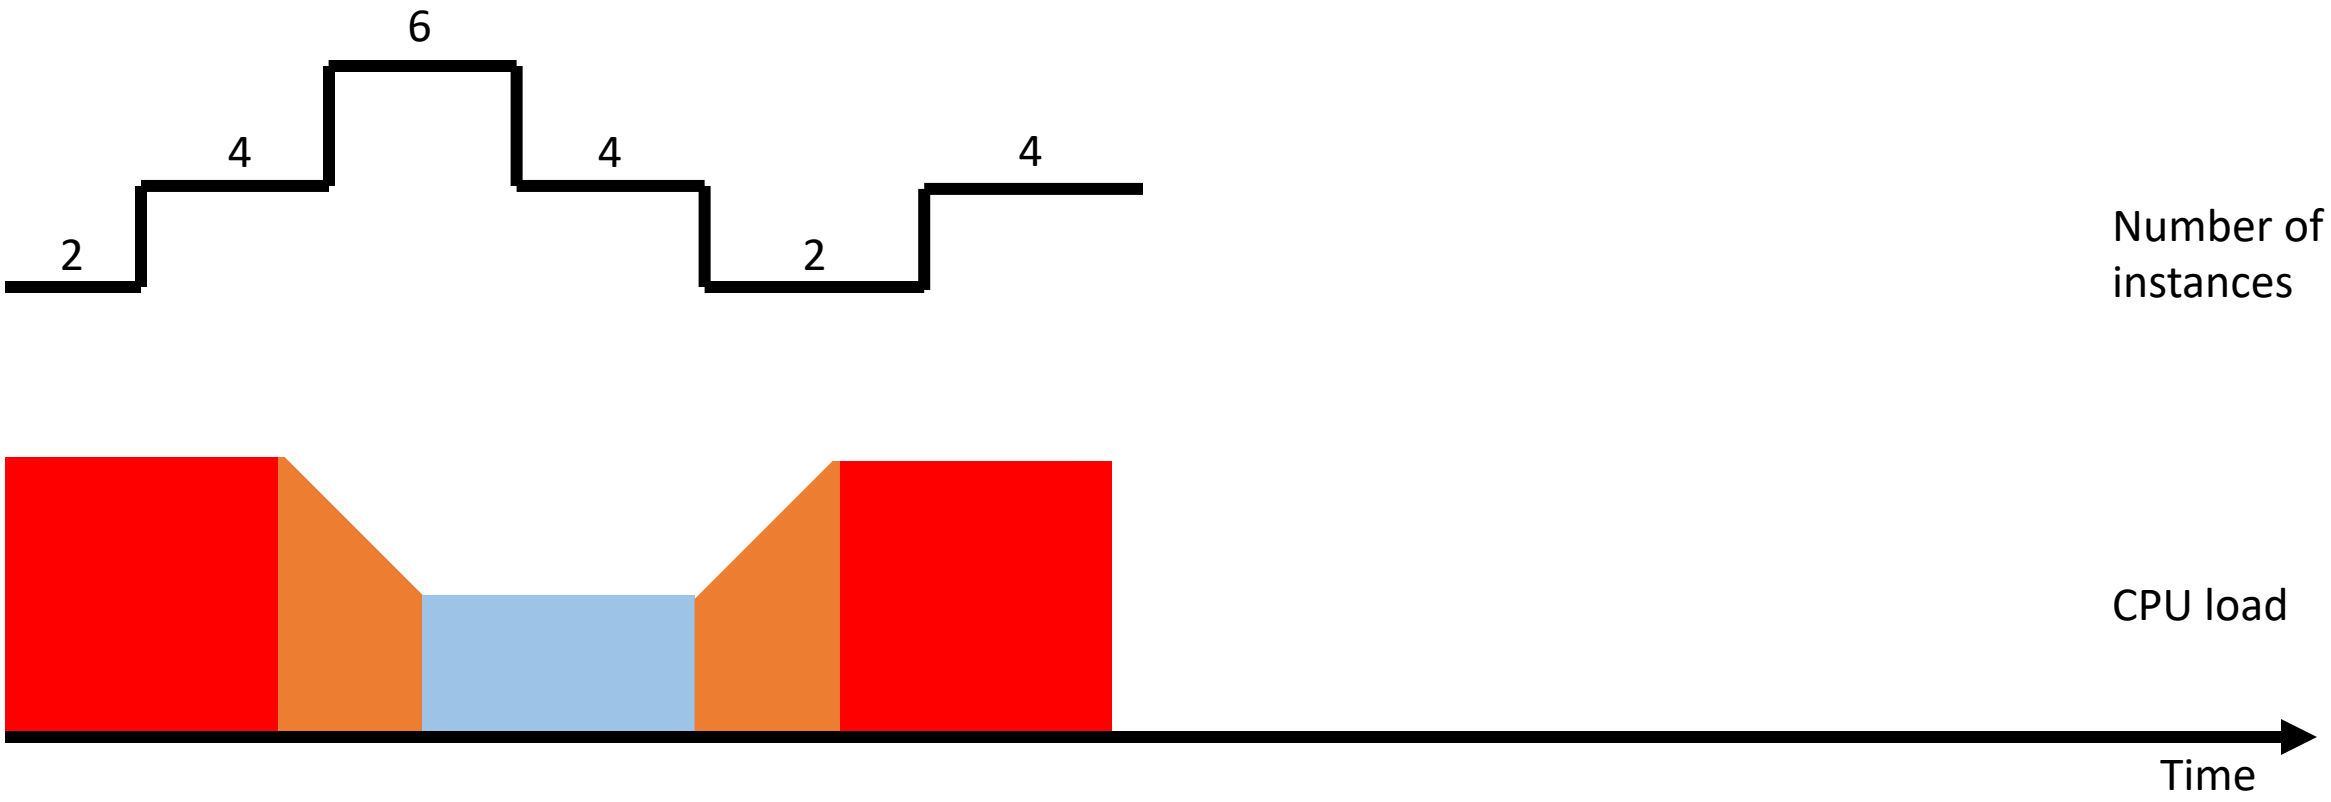

Hysteresis may lead to instability

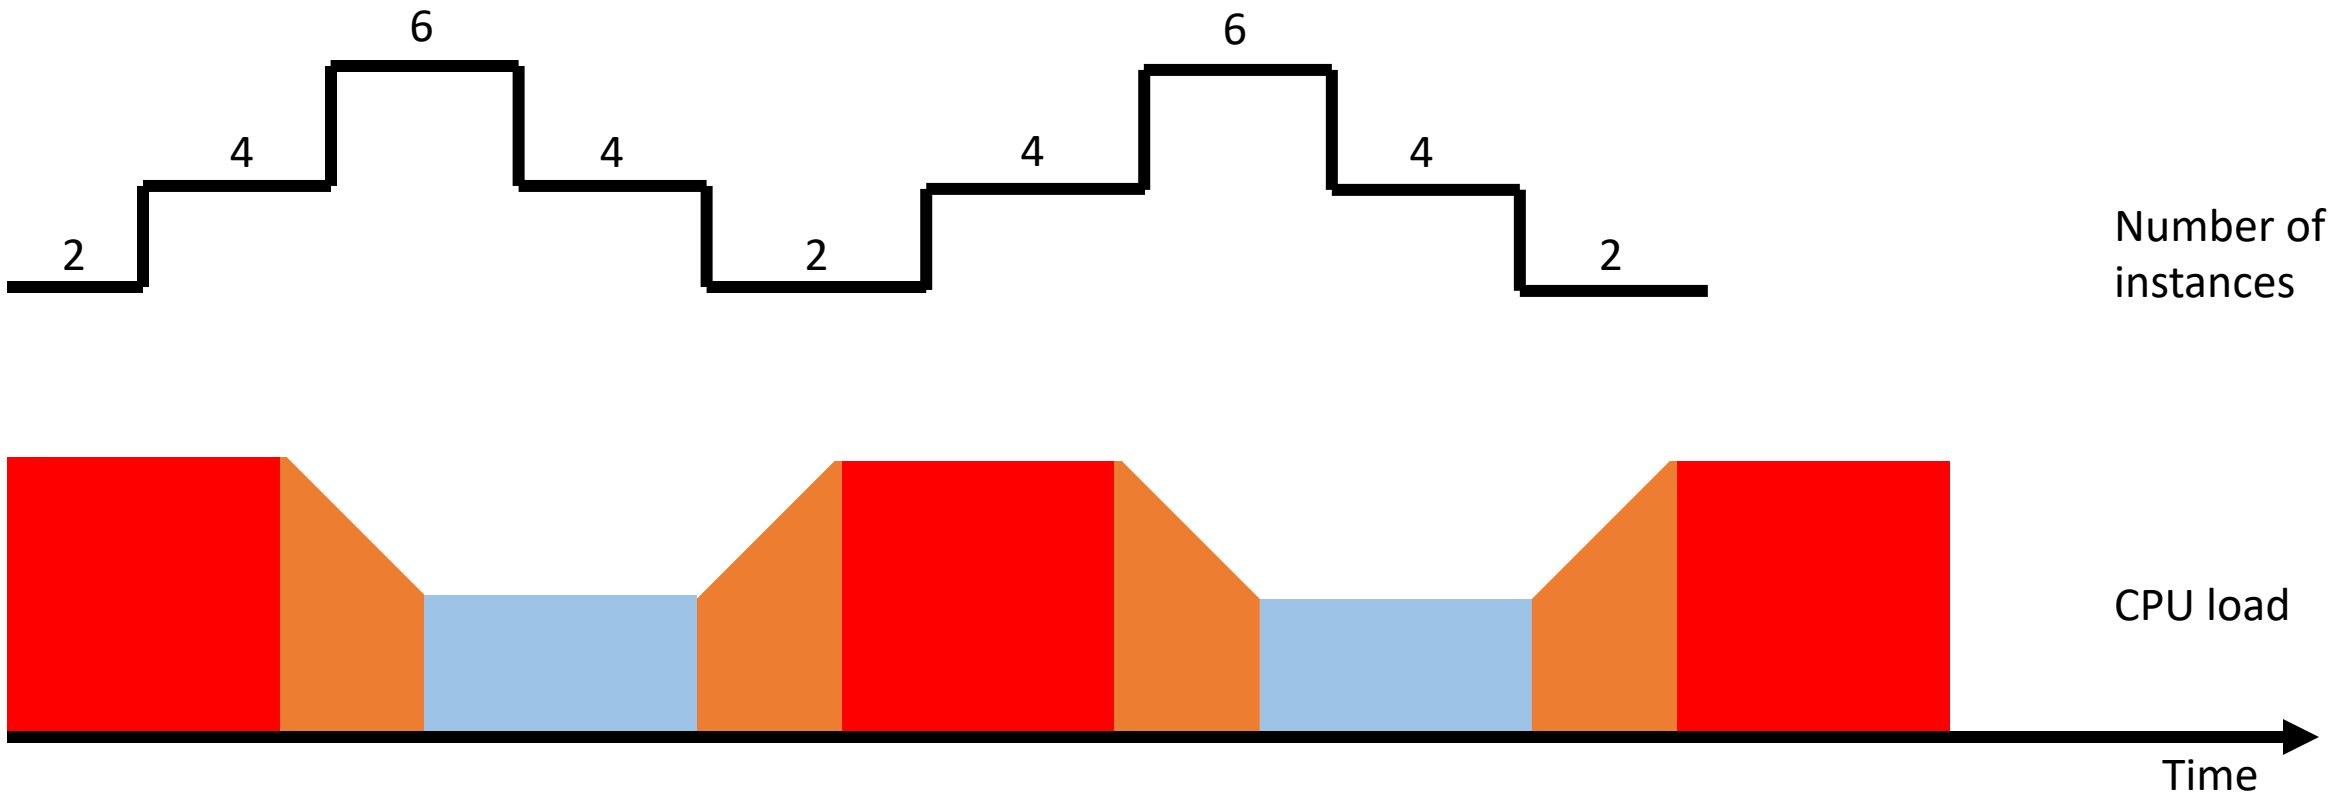

Instability vs dynamics

- This example consider that the workload is constant

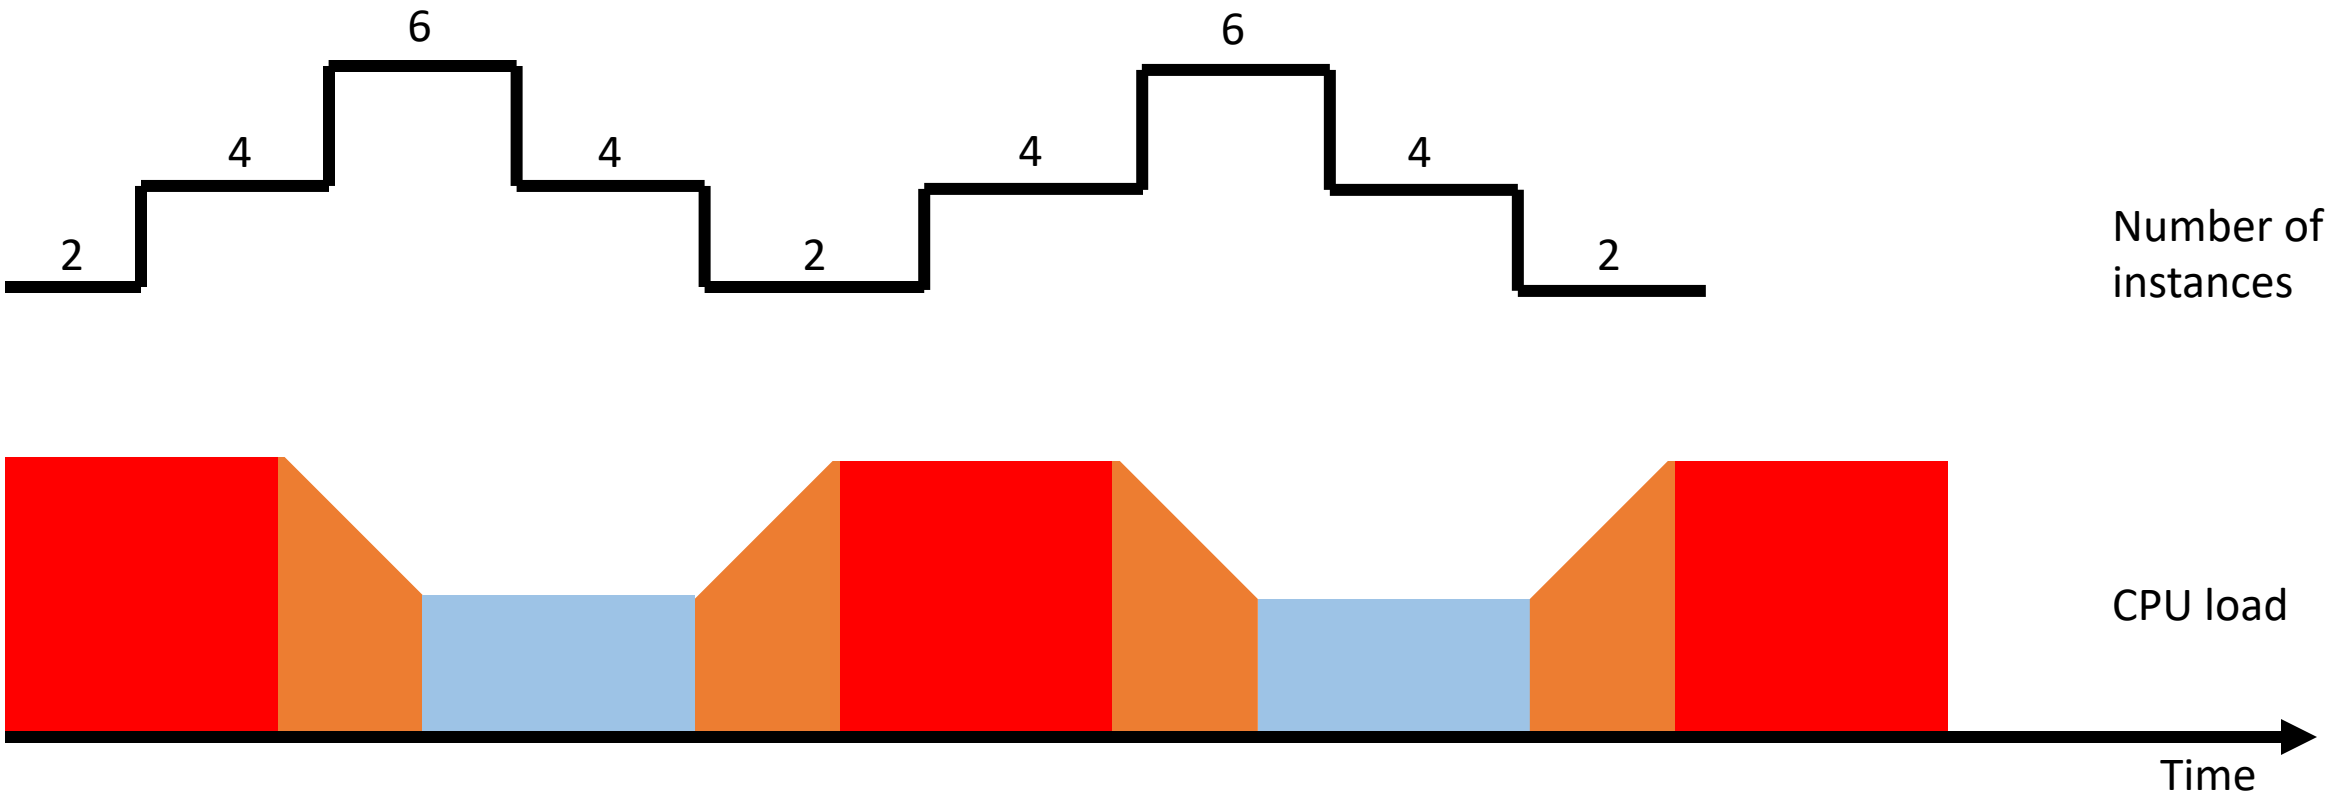

Instability vs dynamics

- If the workload is changing, a system may never stabilize

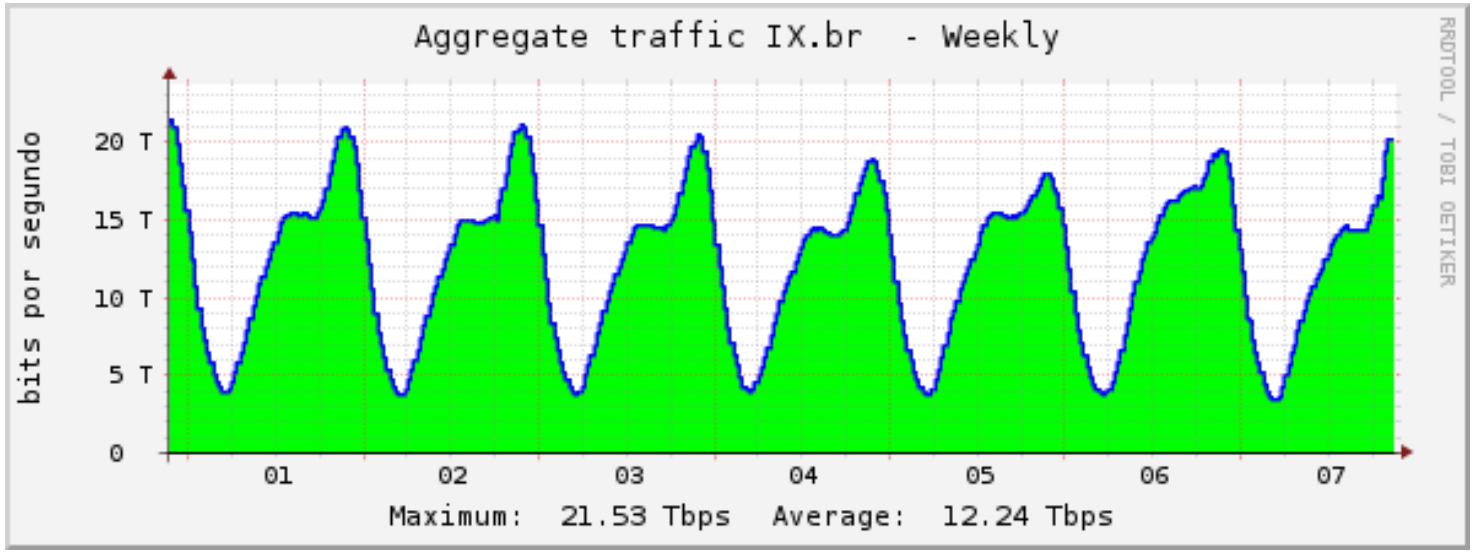

Bits per second at an Internet exchange point in Brazil

Action intensity

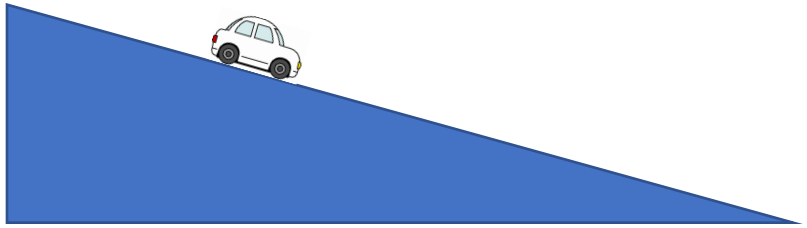

Cruise control at 60km/h

Breaking super hard, releasing the brake, then breaking super hard again…
Would be very uncomfortable and possibly make people sick!

Response intensity

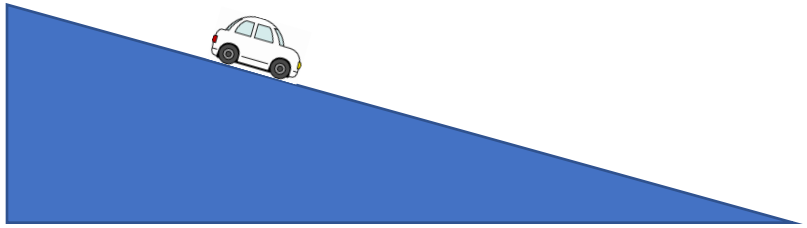

Cruise control at 60km/h

Brake slightly

Response intensity

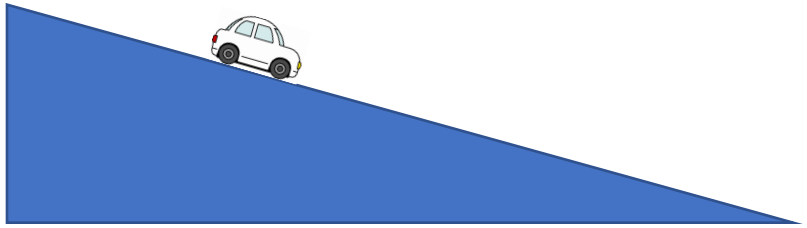

Cruise control at 60km/h
Brake slightly

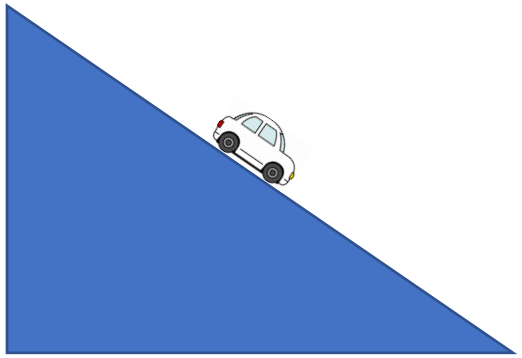

Cruise control at 60km/h
Brake harder

Why scale-out by 2 instances?!

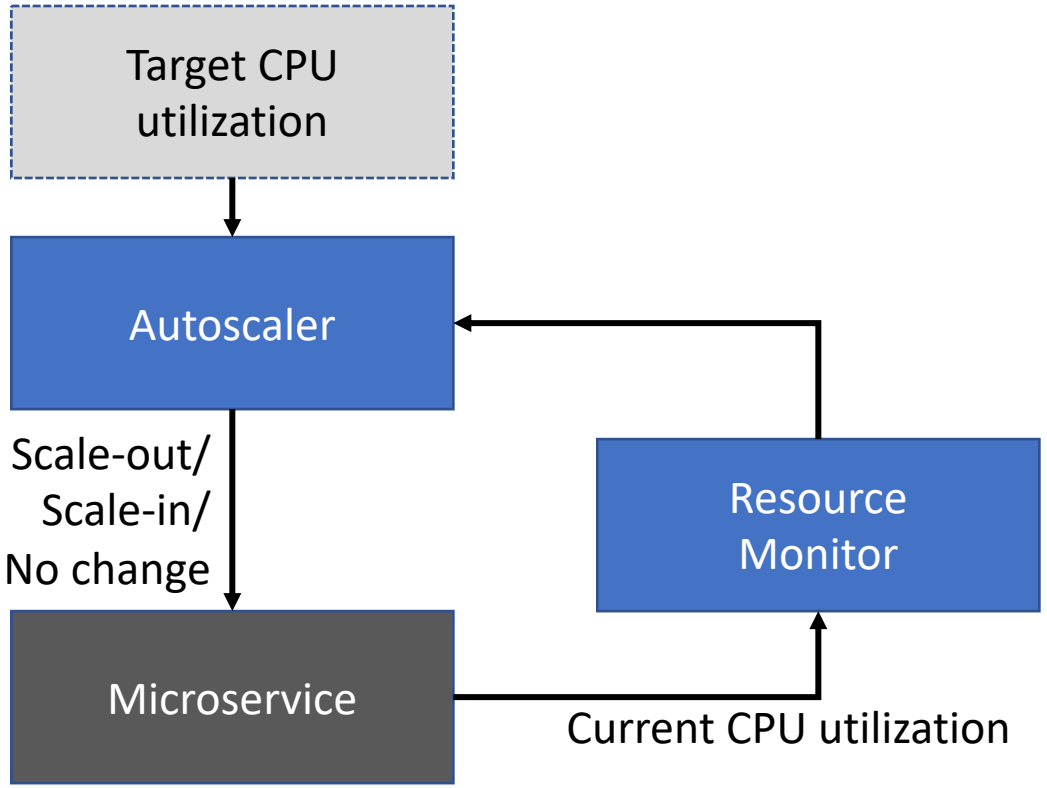

Sequence of events

- T1: Resource monitor tells autoscaler that the microservice is overloaded
- T2: Scale-out **2** instances
- T3: New instances start
- T4: Load balancer starts to distribute load
- T5: Resource monitor tells autoscaler that the microservice is overloaded
- T6: Scale-out **2** instances

CompSci 401: Cloud Computing

Kubernetes Controllers

Prof. Ítalo Cunha

昆山杜克大学
DUKE KUNSHAN UNIVERSITY

Kubernetes controllers

- Run a control loop indefinitely
- Compare actual and desired state of the system
- Make adjustments to achieve the desired state

Event-based control

- Monitors report events to controller
- No need to poll resources periodically for their state
- Only act on state changes

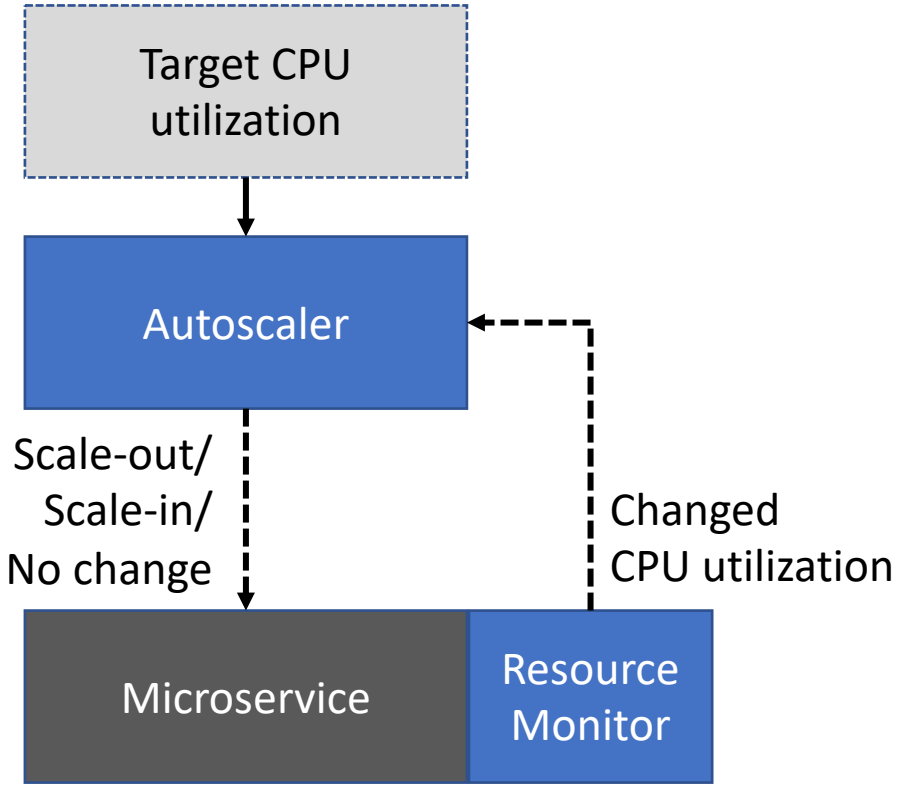

Event-based control

- Monitors report events to controller
- No need to poll resources periodically for their state
- Only act on state changes

- **Reduced overhead and quicker responses**

Example Kubernetes controllers

- Node controller
  - Responsible for noticing and responding when nodes fail

Example Kubernetes controllers

- Node controller
  - Responsible for noticing and responding when nodes fail

```yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.14.2
        ports:
        - containerPort: 80
```

Example Kubernetes controllers

- Node controller
  - Responsible for noticing and responding when nodes go down
- Job controller
  - Watches for Job objects that represent one-off tasks
  - Then creates Pods to run those tasks to completion
- Endpoints controller
  - Naming and discovery to integrate services and jobs
- Service account & token controllers
  - Create default accounts and API access tokens

Example Kubernetes controllers

- Node controller
  - Responsible for noticing and responding when nodes go down
- Job controller
  - Watches for Job objects that represent one-off tasks
  - Then creates Pods to run those tasks to completion.
- Endpoints controller
  - Naming and discovery to integrate services and jobs
- Service account & token controllers
  - Create default accounts and API access tokens
- … we can define or own Resources and Controllers

Kubernetes cluster

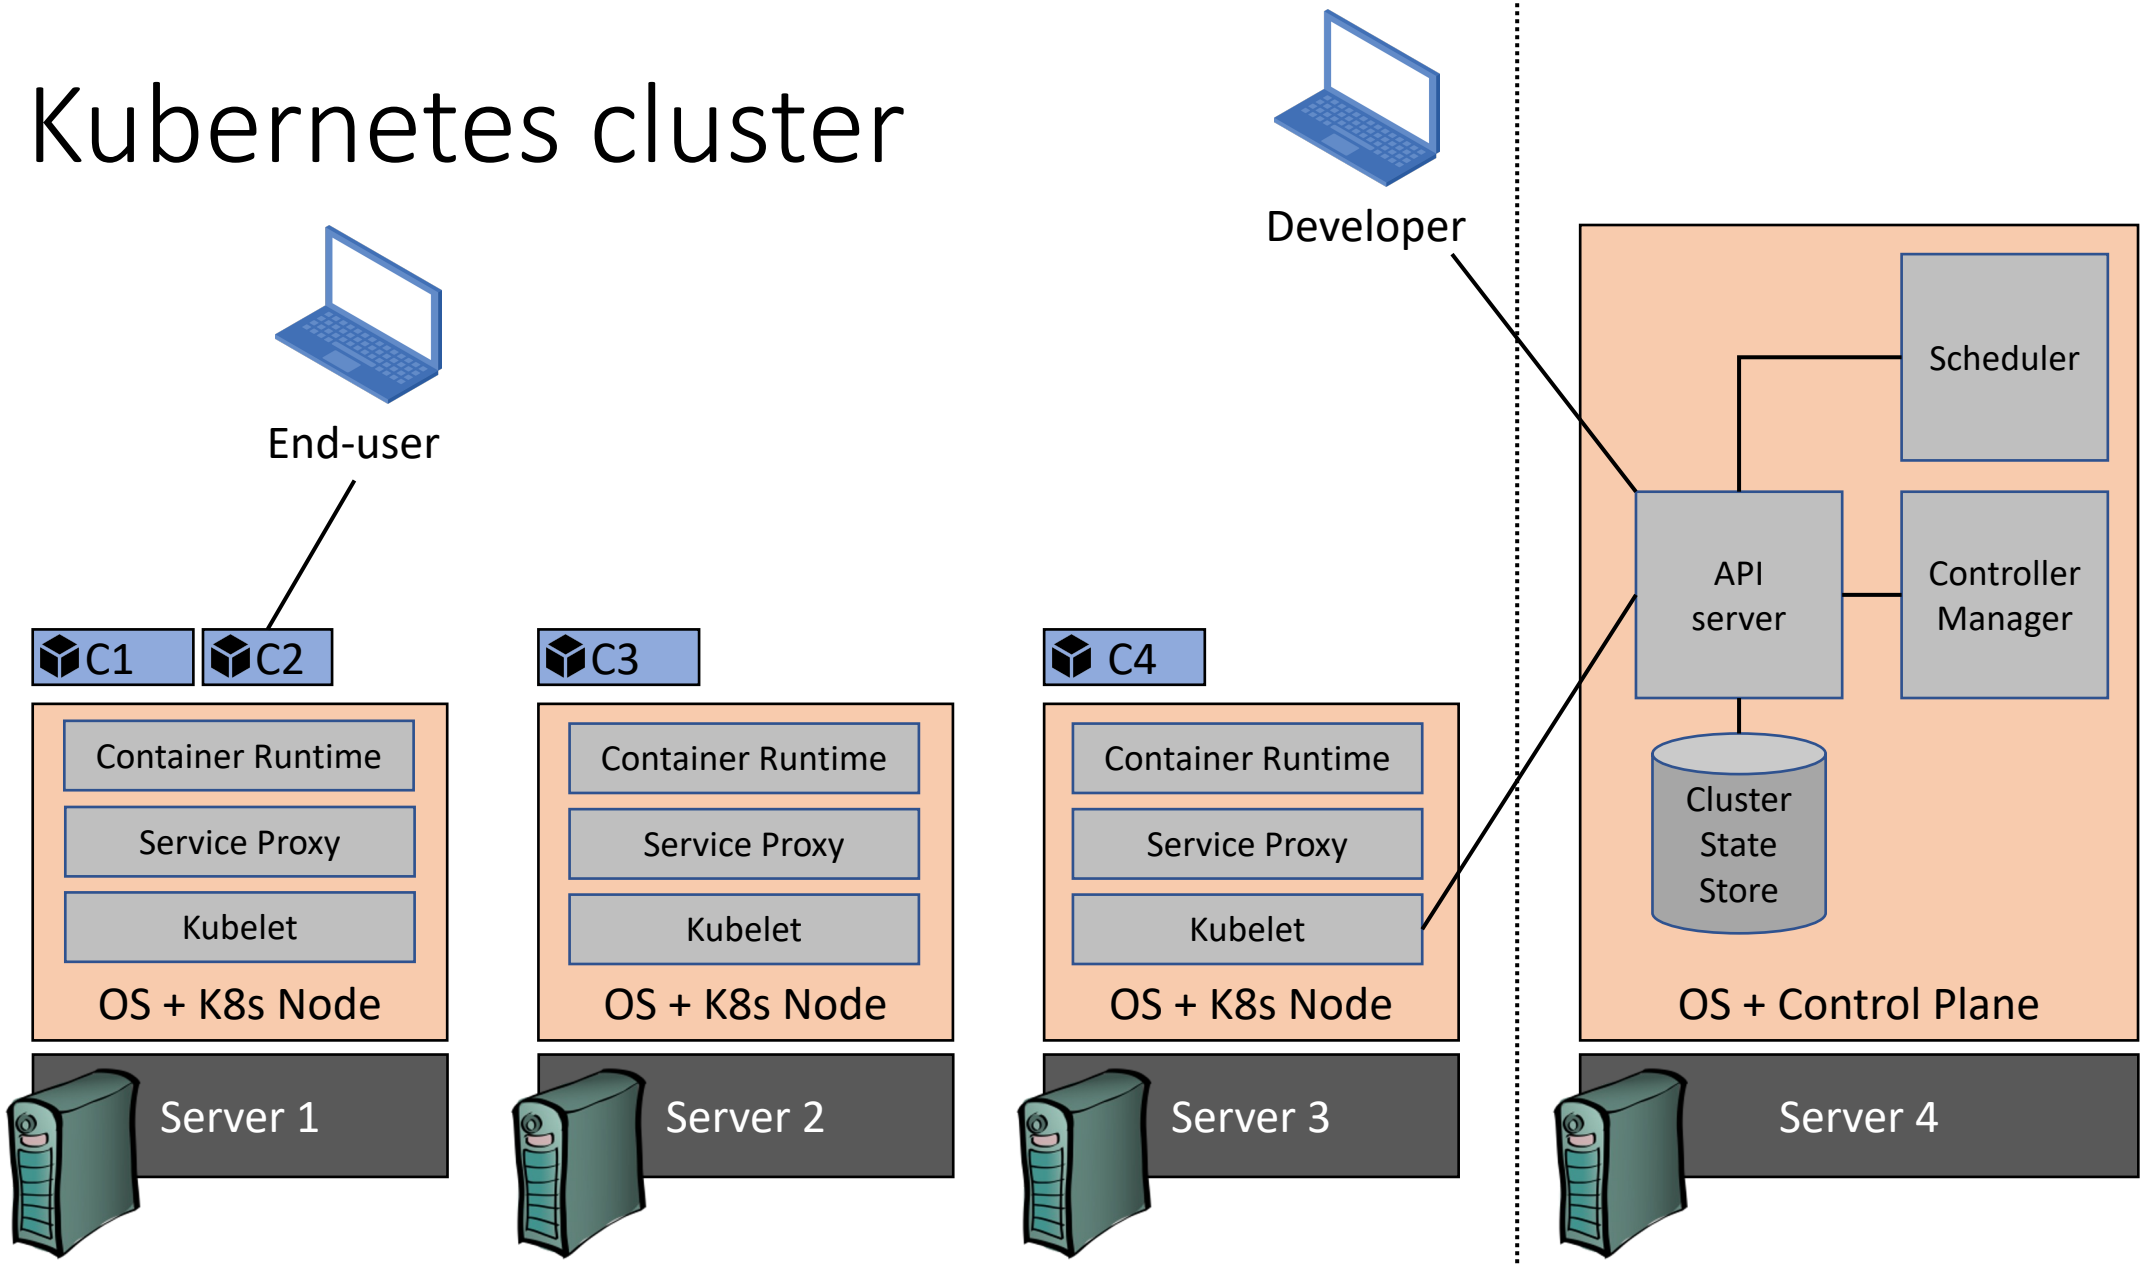

Autoscaling controller

- Control loop delay and hysteresis
  - Stabilization window
  - Period

```yaml
behavior:
  scaleDown:
    stabilizationWindowSeconds: 300
    policies:
    - type: Percent
      value: 100
      periodSeconds: 15
  scaleUp:
    stabilizationWindowSeconds: 0
    policies:
    - type: Percent
      value: 100
      periodSeconds: 15
    - type: Pods
      value: 4
      periodSeconds: 15
    selectPolicy: Max
```

Autoscaling controller

- Control loop delay and hysteresis
  - Stabilization window
  - Period

```yaml
behavior:
  scaleDown:
    stabilizationWindowSeconds: 300
    policies:
    - type: Percent
      value: 100
      periodSeconds: 15
  scaleUp:
    stabilizationWindowSeconds: 0
    policies:
    - type: Percent
      value: 100
      periodSeconds: 15
    - type: Pods
      value: 4
      periodSeconds: 15
    selectPolicy: Max
```

Autoscaling controller

- Control loop delay and hysteresis
  - Stabilization window
  - Period
- Action intensity

$$\text{Instances}_{\text{new}} = \left\lceil \text{Instances}_{\text{old}} \left(\frac{\text{currentState}}{\text{desiredState}} \right) \right\rceil$$

```yaml
behavior:
  scaleDown:
    stabilizationWindowSeconds: 300
    policies:
    - type: Percent
      value: 100
      periodSeconds: 15
  scaleUp:
    stabilizationWindowSeconds: 0
    policies:
    - type: Percent
      value: 100
      periodSeconds: 15
    - type: Pods
      value: 4
      periodSeconds: 15
    selectPolicy: Max
```